

Chapter 2. Input and Output

In this chapter

- [2.1 Input/Output Streams](#)
 - [2.2 Text Input and Output](#)
 - [2.3 Reading and Writing Binary Data](#)
 - [2.4 Object Input/Output Streams and Serialization](#)
 - [2.5 Working with Files](#)
 - [2.6 Memory-Mapped Files](#)
 - [2.7 Regular Expressions](#)
-

In this chapter, we will cover the Java Application Programming Interfaces (APIs) for input and output. You will learn how to access files and directories and how to read and write data in binary and text format. This chapter also shows you the object serialization mechanism that lets you store objects as easily as you can store text or numeric data. Next, we will turn to working with files and directories. We finish the chapter with a discussion of regular expressions, even though they are not actually related to input and output. We couldn't find a better place to handle that topic, and apparently neither could the Java team—the regular expression API specification was attached to a specification request for “new I/O” features.

2.1 Input/Output Streams

In the Java API, an object from which we can read a sequence of bytes is called an *input stream*. An object to which we can write a sequence of bytes is called an *output stream*. These sources and destinations of byte sequences can be—and often are—files, but they can also be network connections and even blocks of memory. The abstract classes `InputStream` and `OutputStream` form the basis for a hierarchy of input/output (I/O) classes.



Note

These input/output streams are unrelated to the streams that you saw in the preceding chapter. For clarity, we will use the terms input stream, output stream, or input/output stream whenever we discuss streams that are used for input and output.

Byte-oriented input/output streams are inconvenient for processing information stored in Unicode (recall that Unicode uses multiple bytes per character). Therefore, a separate hierarchy provides classes for processing Unicode characters that inherit from the abstract `Reader` and `Writer` classes. These classes have read and write operations that are based on two-byte `char` values (that is, UTF-16 code units) rather than `byte` values.

2.1.1 Reading and Writing Bytes

The `InputStream` class has an abstract method:

```
abstract int read()
```

This method reads one byte and returns the byte that was read, or `-1` if it encounters the end of the input source. The designer of a concrete input stream class overrides this method to provide useful functionality. For example, in the `FileInputStream` class, this method reads one byte from a file. `System.in` is a predefined object of a subclass of `InputStream` that allows you to read information from “standard input,” that is, the console or a redirected file.

The `InputStream` class also has nonabstract methods to read an array of bytes or to skip a number of bytes. These methods call the abstract `read` method, so subclasses need to override only one method.

Similarly, the `OutputStream` class defines the abstract method

```
abstract void write(int b)
```

which writes one byte to an output location.

Both the `read` and `write` methods *block* until the byte is actually read or written. This means that if the input stream cannot immediately be accessed (usually because of a busy network connection), the current thread blocks. This gives other threads the chance to do useful work while the method is waiting for the input stream to become available again.

The `available` method lets you check the number of bytes that are currently available for reading. This means a fragment like the following is unlikely to block:

[Click here to view code image](#)

```
int bytesAvailable = in.available();
if (bytesAvailable > 0)
{
    byte[] data = new byte[bytesAvailable];
    in.read(data);
}
```

When you have finished reading or writing to an input/output stream, close it by calling the `close` method. This call frees up the operating system resources that are in limited supply. If an application opens too many input/output streams without closing them, system resources can become depleted. Closing an output stream also *flushes* the buffer used for the output stream: Any bytes that were temporarily placed in a buffer so that they could be delivered as a larger packet are sent off. In particular, if you do not close a file, the last packet of bytes might never be delivered. You can also manually flush the output with the `flush` method.

Even if an input/output stream class provides concrete methods to work with the raw `read` and `write` functions, application programmers rarely use them. The data that you are interested in probably contain numbers, strings, and objects, not raw bytes.

Instead of working with bytes, you can use one of many input/output stream classes derived from the basic `InputStream` and `OutputStream` classes.

java.io.InputStream 1.0

- `abstract int read()`

reads a byte of data and returns the byte read; returns -1 at the end of the input stream.

- `int read(byte[] b)`

reads into an array of bytes and returns the actual number of bytes read, or -1 at the end of the input stream; this method reads at most `b.length` bytes.

- `int read(byte[] b, int off, int len)`

reads into an array of bytes and returns the actual number of bytes read, or -1 at the end of the input stream.

Parameters:

<code>b</code>	The array from which to write the data
<code>off</code>	The offset into <code>b</code> to the first byte that will be written
<code>len</code>	The number of bytes to write

- `long skip(long n)`

skips `n` bytes in the input stream, returns the actual number of bytes skipped (which may be less than `n` if the end of the input stream was encountered).

- `int available()`

returns the number of bytes available, without blocking (recall that blocking means that the current thread loses its turn).

- `void close()`

closes the input stream.

- `void mark(int readlimit)`

puts a marker at the current position in the input stream (not all streams support this feature). If more than `readlimit` bytes have been read from the input stream, the stream is allowed to forget the marker.

- `void reset()`

returns to the last marker. Subsequent calls to `read` reread the bytes. If there is no current marker, the input stream is not reset.

- `boolean markSupported()`

returns `true` if the input stream supports marking.

java.io.OutputStream 1.0

- `abstract void write(int n)`

writes a byte of data.

- `void write(byte[] b)`
- `void write(byte[] b, int off, int len)`

writes all bytes or a range of bytes in the array `b`.

Parameters:

<code>b</code>	The array from which to write the data
<code>off</code>	The offset into <code>b</code> to the first byte that will be written
<code>len</code>	The number of bytes to write

- `void close()`

flushes and closes the output stream.

- `void flush()`

flushes the output stream—that is, sends any buffered data to its destination.

2.1.2 The Complete Stream Zoo

Unlike C, which gets by just fine with a single type `FILE*`, Java has a whole zoo of more than 60 (!) different input/output stream types (see [Figures 2.1](#) and [2.2](#)).

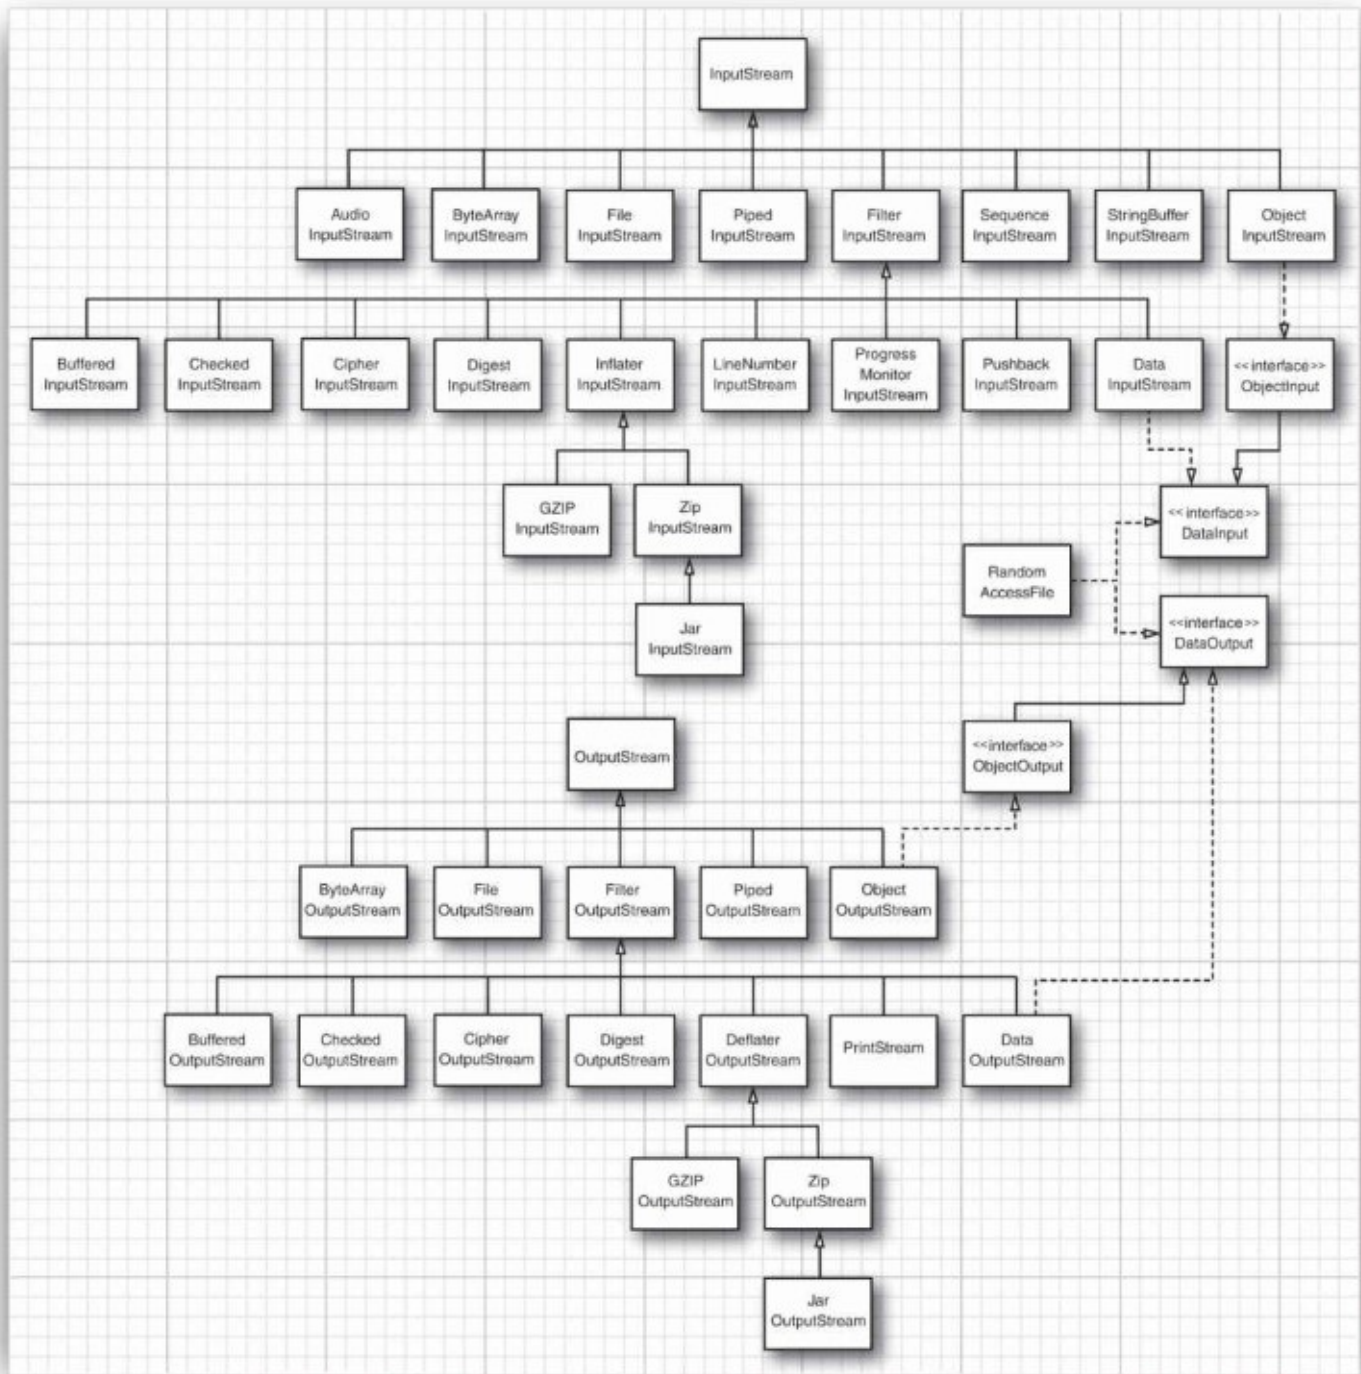


Figure 2.1 Input and output stream hierarchy

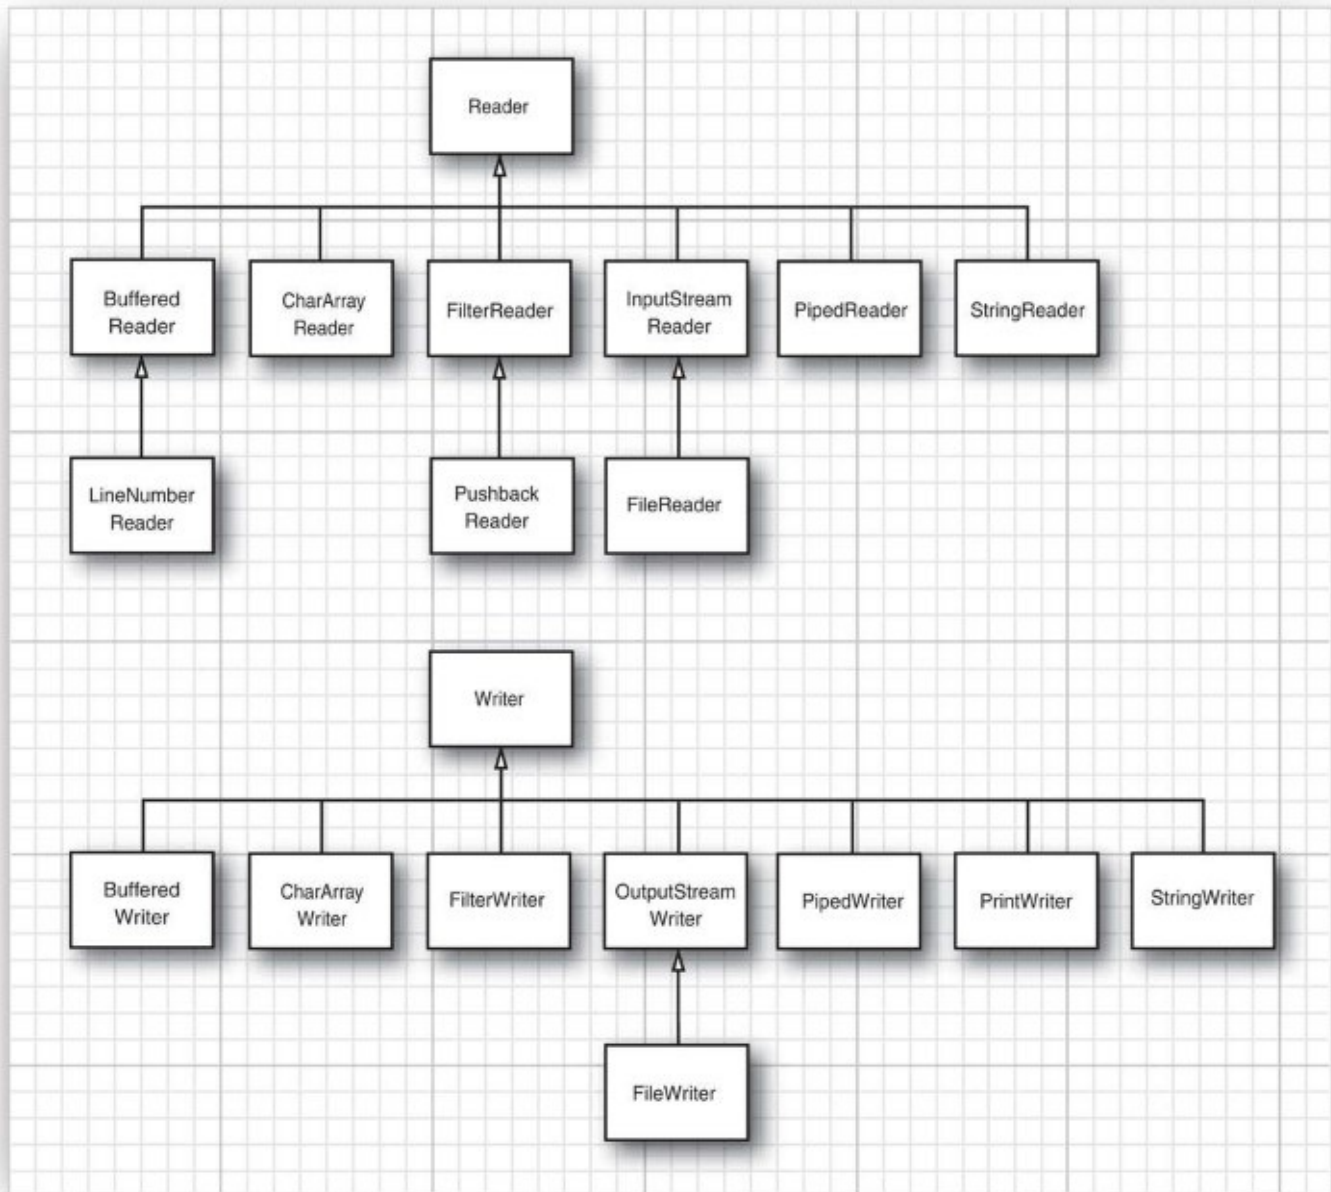


Figure 2.2 Reader and writer hierarchy

Let's divide the animals in the input/output stream zoo by how they are used. There are separate hierarchies for classes that process bytes and characters. As you saw, the `InputStream` and `OutputStream` classes let you read and write individual bytes and arrays of bytes. These classes form the basis of the hierarchy shown in [Figure 2.1](#). To read and write strings and numbers, you need more capable subclasses. For example, `DataInputStream` and `DataOutputStream` let you read and write all the primitive Java types in binary format. Finally, there are input/output streams that do useful stuff; for example, the `ZipInputStream` and `ZipOutputStream` let you read and write files in the familiar ZIP compression format.

For Unicode text, on the other hand, you can use subclasses of the abstract classes `Reader` and `Writer` (see [Figure 2.2](#)). The basic methods of the `Reader` and `Writer` classes are similar to those of `InputStream` and `OutputStream`.

[Click here to view code image](#)


```

abstract int read()
abstract void write(int c)

```

The `read` method returns either a UTF-16 code unit (as an integer between 0 and 65535) or -1 when you have reached the end of the file. The `write` method is called with a Unicode code unit. (See Volume I, [Chapter 3](#) for a discussion of Unicode code units.)

There are four additional interfaces: `Closeable`, `Flushable`, `Readable`, and `Appendable` (see [Figure 2.3](#)). The first two interfaces are very simple, with methods

[Click here to view code image](#)

```

void close() throws IOException

```

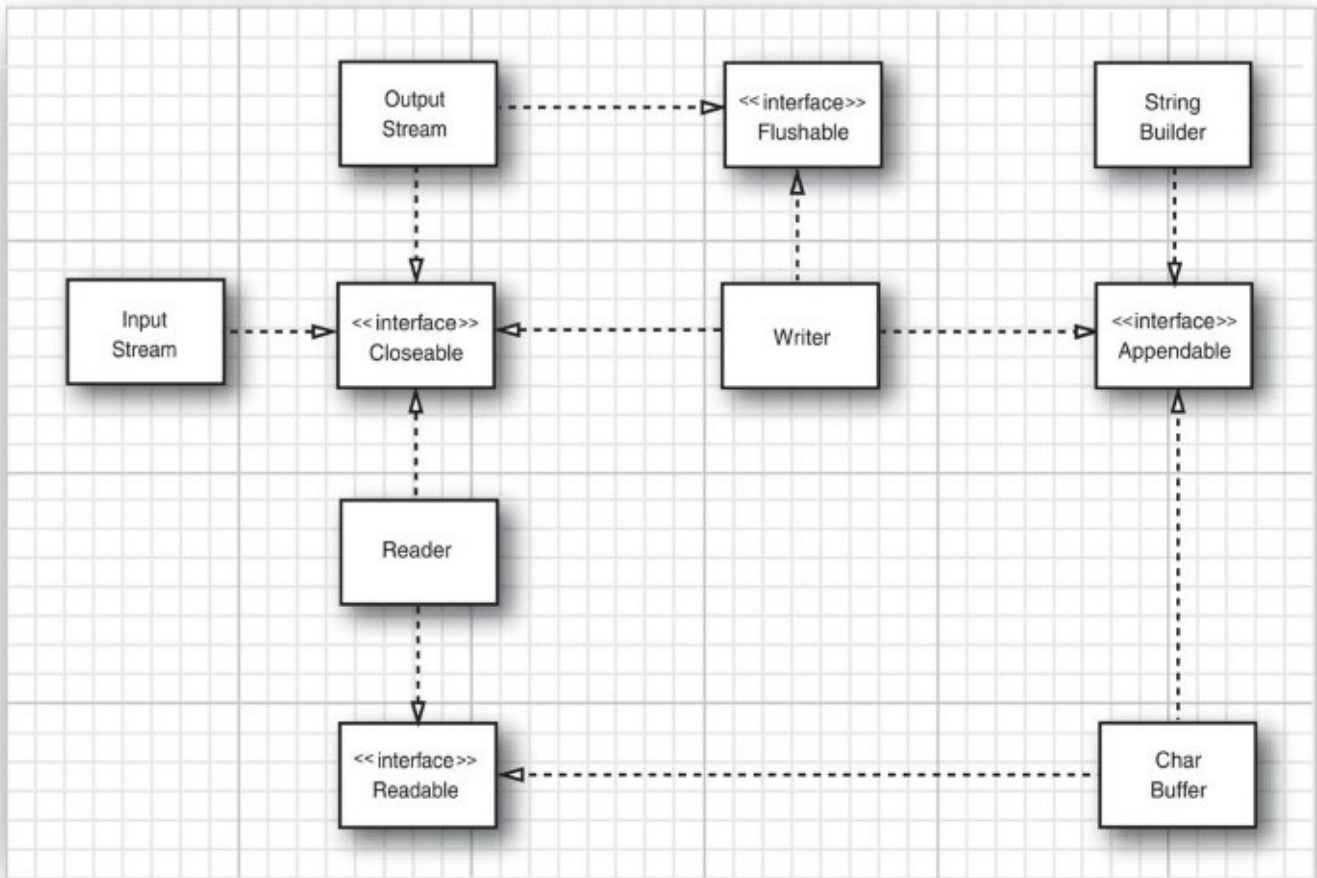


Figure 2.3 The `Closeable`, `Flushable`, `Readable`, and `Appendable` interfaces

and

```

void flush()

```

respectively. The classes `InputStream`, `OutputStream`, `Reader`, and `Writer` all implement the `Closeable` interface.



Note

The `java.io.Closeable` interface extends the `java.lang.AutoCloseable` interface. Therefore, you can use the `try-with-resources` statement with any `Closeable`. Why have two interfaces? The `close` method of the `Closeable` interface only throws an `IOException`, whereas the `AutoCloseable.close` method may throw any exception.

`OutputStream` and `Writer` implement the `Flushable` interface.

The `Readable` interface has a single method

```
int read(CharBuffer cb)
```

The `CharBuffer` class has methods for sequential and random read/write access. It represents an in-memory buffer or a memory-mapped file. (See [Section 2.6.2](#), “[The Buffer Data Structure](#),” on [p. 124](#) for details.)

The `Appendable` interface has two methods for appending single characters and character sequences:

[Click here to view code image](#)

```
Appendable append(char c)
Appendable append(CharSequence s)
```

The `CharSequence` interface describes basic properties of a sequence of `char` values. It is implemented by `String`, `CharBuffer`, `StringBuilder`, and `StringBuffer`.

Of the input/output stream classes, only `Writer` implements `Appendable`.

java.io.Closeable 5.0

- `void close()`

closes this `Closeable`. This method may throw an `IOException`.

java.io.Flushable 5.0

- `void flush()`

flushes this `Flushable`.

java.lang.Readable 5.0

- `int read(CharBuffer cb)`

attempts to read as many `char` values into `cb` as it can hold. Returns the number of values read, or `-1` if no further values are available from this `Readable`.

***java.lang.Appendable* 5.0**

- `Appendable append(char c)`
- `Appendable append(CharSequence cs)`

appends the given code unit, or all code units in the given sequence, to this `Appendable`; returns this.

***java.lang.CharSequence* 1.4**

- `char charAt(int index)`
returns the code unit at the given index.
 - `int length()`
returns the number of code units in this sequence.
 - `CharSequence subSequence(int startIndex, int endIndex)`
returns a `CharSequence` consisting of the code units stored from index `startIndex` to `endIndex - 1`.
 - `String toString()`
returns a string consisting of the code units of this sequence.
-

2.1.3 Combining Input/Output Stream Filters

`FileInputStream` and `FileOutputStream` give you input and output streams attached to a disk file. You need to pass the file name or full path name of the file to the constructor. For example,

[Click here to view code image](#)

```
FileInputStream fin = new FileInputStream("employee.dat");
```

looks in the user directory for a file named `employee.dat`.



Tip

All the classes in `java.io` interpret relative path names as starting from the user's working directory. You can get this directory by a call to `System.getProperty("user.dir")`.



Caution

Since the backslash character is the escape character in Java strings, be sure to use `\\` for Windows-style path names (for example, `C:\\Windows\\win.ini`). In Windows, you can also use a single forward slash (`C:/Windows/win.ini`) because most Windows file-handling system calls will interpret forward slashes as file separators. However, this is not recommended—the behavior of the Windows system functions is subject to change. Instead, for portable programs, use the file separator character for the platform on which your program runs. It is available as the constant string `java.io.File.separator`.

Like the abstract `InputStream` and `OutputStream` classes, these classes only support reading and writing at the byte level. That is, we can only read bytes and byte arrays from the object `fin`.

```
byte b = (byte) fin.read();
```

As you will see in the next section, if we just had a `DataInputStream`, we could read numeric types:

[Click here to view code image](#)

```
DataInputStream din = . . . ;  
double x = din.readDouble();
```

But just as the `FileInputStream` has no methods to read numeric types, the `DataInputStream` has no method to get data from a file.

Java uses a clever mechanism to separate two kinds of responsibilities. Some input streams (such as the `FileInputStream` and the input stream returned by the `openStream` method of the `URL` class) can retrieve bytes from files and other more exotic locations. Other input streams (such as the `DataInputStream`) can assemble bytes into more useful data types. The Java programmer has to combine the two. For example, to be able to read numbers from a file, first create a `FileInputStream` and then pass it to the constructor of a `DataInputStream`.

[Click here to view code image](#)

```
FileInputStream fin = new FileInputStream("employee.dat");  
DataInputStream din = new DataInputStream(fin);  
double x = din.readDouble();
```

If you look at [Figure 2.1](#) again, you can see the classes `FilterInputStream` and `FilterOutputStream`. The subclasses of these classes are used to add capabilities to input/output streams that process bytes.

You can add multiple capabilities by nesting the filters. For example, by default, input streams are not buffered. That is, every call to `read` asks the operating system to dole out yet another byte. It is more efficient to request blocks of data instead and store them in a buffer. If you want buffering *and* the data input methods for a file, you need to use the following rather monstrous sequence of constructors:

[Click here to view code image](#)

```
DataInputStream din = new DataInputStream(  
    new BufferedInputStream(  
        new FileInputStream(  
            new File("employee.dat")  
        )  
    )  
);
```

```
new FileInputStream("employee.dat")));
```

Notice that we put the `DataInputStream` *last* in the chain of constructors because we want to use the `DataInputStream` methods, and we want *them* to use the buffered `read` method.

Sometimes you'll need to keep track of the intermediate input streams when chaining them together. For example, when reading input, you often need to peek at the next byte to see if it is the value that you expect. Java provides the `PushbackInputStream` for this purpose.

[Click here to view code image](#)

```
PushbackInputStream pbin = new PushbackInputStream(  
    new BufferedInputStream(  
        new FileInputStream("employee.dat")));
```

Now you can speculatively read the next byte

```
int b = pbin.read();
```

and throw it back if it isn't what you wanted.

```
if (b != '<') pbin.unread(b);
```

However, reading and unreading are the *only* methods that apply to a pushback input stream. If you want to look ahead and also read numbers, then you need both a pushback input stream and a data input stream reference.

[Click here to view code image](#)

```
DataInputStream din = new DataInputStream(  
    pbin = new PushbackInputStream(  
        new BufferedInputStream(  
            new FileInputStream("employee.dat"))));
```

Of course, in the input/output libraries of other programming languages, niceties such as buffering and lookahead are automatically taken care of, so it is a bit of a hassle to resort, in Java, to combining stream filters. However, the ability to mix and match filter classes to construct truly useful sequences of input/output streams does give you an immense amount of flexibility. For example, you can read numbers from a compressed ZIP file by using the following sequence of input streams (see [Figure 2.4](#)):

[Click here to view code image](#)

```
ZipInputStream zin = new ZipInputStream(new FileInputStream("employee.zip"));  
DataInputStream din = new DataInputStream(zin);
```

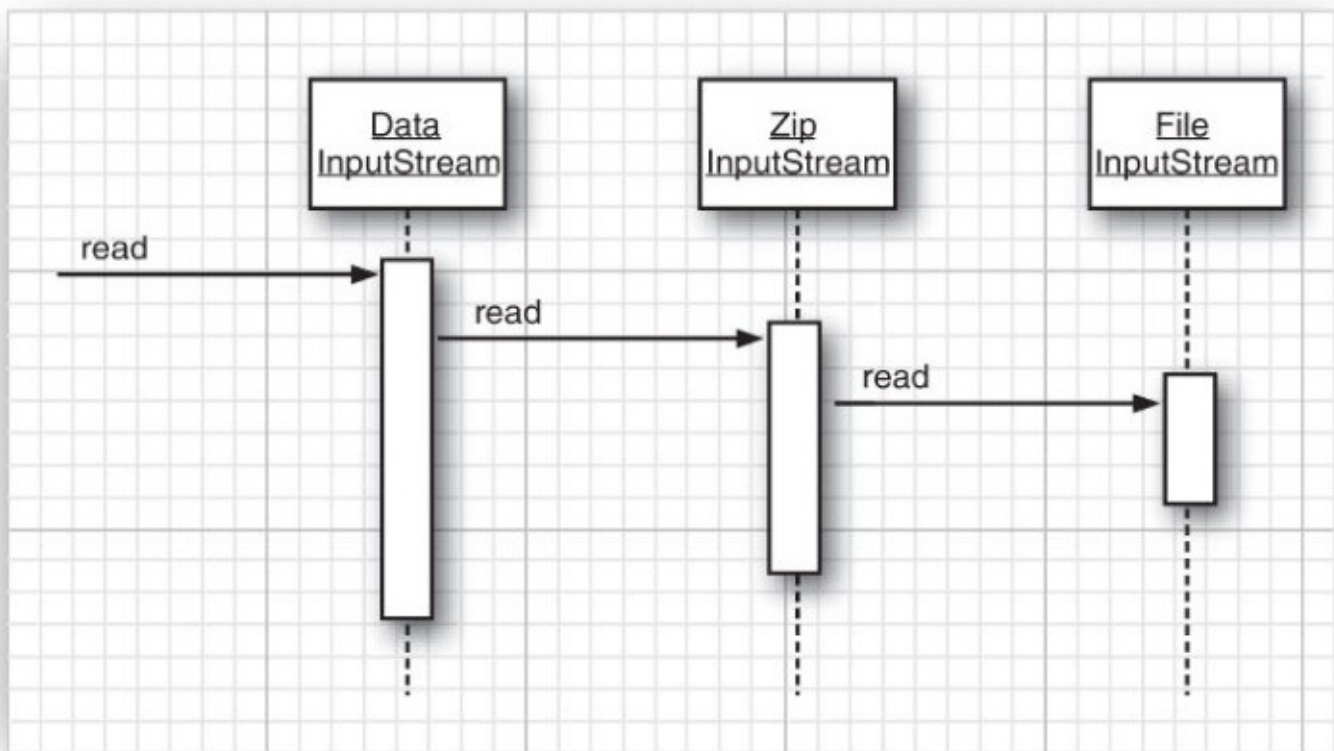


Figure 2.4 A sequence of filtered input streams

(See [Section 2.3.3](#), “[ZIP Archives](#),” on [p. 77](#) for more on Java’s handling of ZIP files.)

java.io.FileInputStream 1.0

- `FileInputStream(String name)`
- `FileInputStream(File file)`

creates a new file input stream using the file whose path name is specified by the `name` string or the `file` object. (The `File` class is described at the end of this chapter.) Path names that are not absolute are resolved relative to the working directory that was set when the VM started.

java.io.FileOutputStream 1.0

- `FileOutputStream(String name)`
- `FileOutputStream(String name, boolean append)`
- `FileOutputStream(File file)`
- `FileOutputStream(File file, boolean append)`

creates a new file output stream specified by the `name` string or the `file` object. (The `File` class is described at the end of this chapter.) If the `append` parameter is `true`, an existing file with the same name will not be deleted and data will be added at the end of the file. Otherwise, this method deletes any existing file with the same name.

`java.io.BufferedInputStream 1.0`

- `BufferedInputStream(InputStream in)`

creates a buffered input stream. A buffered input stream reads bytes from a stream without causing a device access every time. When the buffer is empty, a new block of data is read into the buffer.

`java.io.BufferedOutputStream 1.0`

- `BufferedOutputStream(OutputStream out)`

creates a buffered output stream. A buffered output stream collects bytes to be written without causing a device access every time. When the buffer fills up or when the stream is flushed, the data are written.

`java.io.PushbackInputStream 1.0`

- `PushbackInputStream(InputStream in)`
- `PushbackInputStream(InputStream in, int size)`

constructs an input stream with one-byte lookahead or a pushback buffer of specified size.

- `void unread(int b)`

pushes back a byte, which is retrieved again by the next call to `read`.

Parameters: `b` The byte to be read again

2.2 Text Input and Output

When saving data, you have the choice between binary and text formats. For example, if the integer 1234 is saved in binary, it is written as the sequence of bytes 00 00 04 D2 (in hexadecimal notation). In text format, it is saved as the string "1234". Although binary I/O is fast and efficient, it is not easily readable by humans. We first discuss text I/O and cover binary I/O in [Section 2.3](#), "[Reading and Writing Binary Data](#)," on [p. 69](#).

When saving text strings, you need to consider the *character encoding*. In the UTF-16 encoding that Java uses internally, the string "José" is encoded as 00 4A 00 6F 00 73 00 E9 (in hex). However, many programs expect that text files are encoded in a different encoding. In UTF-8, the encoding most commonly used on the Internet, the string would be written as 4A 6F 73 C3 A9, without the zero bytes for the first three letters and with two bytes for the é character.

The `OutputStreamWriter` class turns an output stream of Unicode code units into a stream of bytes, using a chosen character encoding. Conversely, the `InputStreamReader` class turns an input stream that contains bytes (specifying characters in some character encoding) into a reader that emits Unicode code units.

For example, here is how you make an input reader that reads keystrokes from the console and converts them to Unicode:

[Click here to view code image](#)

```
Reader in = new InputStreamReader(System.in);
```

This input stream reader assumes the default character encoding used by the host system. On desktop operating systems, that can be an archaic encoding such as Windows 1252 or MacRoman. You should always choose a specific encoding in the constructor for the `InputStreamReader`, for example:

[Click here to view code image](#)

```
Reader in = new InputStreamReader(new FileInputStream("data.txt"),  
StandardCharsets.UTF_8);
```

See [Section 2.2.4](#), “[Character Encodings](#),” on [p. 67](#) for more information on character encodings.

2.2.1 How to Write Text Output

For text output, use a `PrintWriter`. That class has methods to print strings and numbers in text format. There is a convenience constructor for printing to a file. The statement

[Click here to view code image](#)

```
PrintWriter out = new PrintWriter("employee.txt", "UTF-8");
```

is equivalent to

[Click here to view code image](#)

```
PrintWriter out = new PrintWriter(  
    new FileOutputStream("employee.txt"), "UTF-8");
```

To write to a print writer, use the same `print`, `println`, and `printf` methods that you used with `System.out`. You can use these methods to print numbers (`int`, `short`, `long`, `float`, `double`), characters, `boolean` values, strings, and objects.

For example, consider this code:

[Click here to view code image](#)

```
String name = "Harry Hacker";  
double salary = 75000;  
out.print(name);  
out.print(' ');  
out.println(salary);
```

This writes the characters

```
Harry Hacker 75000.0
```

to the writer `out`. The characters are then converted to bytes and end up in the file `employee.txt`.

The `println` method adds the correct end-of-line character for the target system ("`\r\n`" on Windows, "`\n`" on UNIX) to the line. This is the string obtained by the call

```
System.getProperty("line.separator").
```

If the writer is set to *autoflush mode*, all characters in the buffer are sent to their destination whenever `println` is called. (Print writers are always buffered.) By default, autoflushing is *not*

enabled. You can enable or disable autoflushing by using the `PrintWriter(Writer writer, boolean autoFlush)` constructor:

[Click here to view code image](#)

```
PrintWriter out = new PrintWriter(  
    new OutputStreamWriter(  
        new FileOutputStream("employee.txt"), "UTF-8"),  
    true); // autoflush
```

The `print` methods don't throw exceptions. You can call the `checkError` method to see if something went wrong with the output stream.



Note

Java veterans might wonder whatever happened to the `PrintStream` class and to `System.out`. In Java 1.0, the `PrintStream` class simply truncated all Unicode characters to ASCII characters by dropping the top byte. (At the time, Unicode was still a 16-bit encoding.) Clearly, that was not a clean or portable approach, and it was fixed with the introduction of readers and writers in Java 1.1. For compatibility with existing code, `System.in`, `System.out`, and `System.err` are still input/output streams, not readers and writers. But now the `PrintStream` class internally converts Unicode characters to the default host encoding in the same way the `PrintWriter` does. Objects of type `PrintStream` act exactly like print writers when you use the `print` and `println` methods, but unlike print writers they allow you to output raw bytes with the `write(int)` and `write(byte[])` methods.

java.io.PrintWriter 1.1

- `PrintWriter(Writer out)`
creates a new `PrintWriter` that writes to the given writer.
 - `PrintWriter(Writer writer)`
creates a new `PrintWriter` that writes to the given file, using the given character encoding.
 - `PrintWriter(String filename, String encoding)`
 - `PrintWriter(File file, String encoding)`
creates a new `PrintWriter` that writes to the given file, using the given character encoding.
 - `void print(Object obj)`
prints an object by printing the string resulting from `toString`.
 - `void print(String s)`
prints a string containing Unicode code units.
 - `void println(String s)`
prints a string followed by a line terminator. Flushes the output stream if it is in autoflush mode.
 - `void print(char[] s)`
prints all Unicode code units in the given array.
 - `void print(char c)`
prints a Unicode code unit.
 - `void print(int i)`
 - `void print(long l)`
 - `void print(float f)`
 - `void print(double d)`
 - `void print(boolean b)`
prints the given value in text format.
 - `void printf(String format, Object... args)`
prints the given values as specified by the format string. See Volume I, [Chapter 3](#) for the specification of the format string.
 - `boolean checkError()`
returns `true` if a formatting or output error occurred. Once the output stream has encountered an error, it is tainted and all calls to `checkError` return `true`.
-

2.2.2 How to Read Text Input

The easiest way to process arbitrary text is the `Scanner` class that we used extensively in Volume I. You can construct a `Scanner` from any input stream.

Alternatively, you can read a short text file into a string like this:

[Click here to view code image](#)

```
String content = new String(Files.readAllBytes(path), charset);
```

But if you want the file as a sequence of lines, call

[Click here to view code image](#)

```
List<String> lines = Files.readAllLines(path, charset);
```

If the file is large, process the lines lazily as a `Stream<String>`:

[Click here to view code image](#)

```
try (Stream<String> lines = Files.lines(path, charset))
{
    . . .
}
```

In early versions of Java, the only game in town for processing text input was the `BufferedReader` class. Its `readLine` method yields a line of text, or `null` when no more input is available. A typical input loop looks like this:

[Click here to view code image](#)

```
InputStream inputStream = . . .;
try (BufferedReader in = new BufferedReader(new InputStreamReader(inputStream,
                                                                    StandardCharsets.
                                                                    UTF_8))
{
    String line;
    while ((line = in.readLine()) != null)
    {
        do something with line
    }
}
```

Nowadays, the `BufferedReader` class also has a `lines` method that yields a `Stream<String>`. However, unlike a `Scanner`, a `BufferedReader` has no methods for reading numbers.

2.2.3 Saving Objects in Text Format

In this section, we walk you through an example program that stores an array of `Employee` records in a text file. Each record is stored in a separate line. Instance fields are separated from each other by delimiters. We use a vertical bar (`|`) as our delimiter. (A colon (`:`) is another popular choice. Part of the fun is that everyone uses a different delimiter.) Naturally, we punt on the issue of what might happen if a `|` actually occurs in one of the strings we save.

Here is a sample set of records:

[Click here to view code image](#)

```
Harry Hacker|35500|1989-10-01
Carl Cracker|75000|1987-12-15
Tony Tester|38000|1990-03-15
```

Writing records is simple. Since we write to a text file, we use the `PrintWriter` class. We simply write all fields, followed by either a `|` or, for the last field, a `\n`. This work is done in the following `writeData` method that we add to our `Employee` class:

[Click here to view code image](#)

```
public static void writeEmployee(PrintWriter out, Employee e)
```

```

{
    out.println(e.getName() + "|" + e.getSalary() + "|" + e.getHireDay());
}

```

To read records, we read in a line at a time and separate the fields. We use a scanner to read each line and then split the line into tokens with the `String.split` method.

[Click here to view code image](#)

```

public static Employee readEmployee(Scanner in)
{
    String line = in.nextLine();
    String[] tokens = line.split("\\|");
    String name = tokens[0];
    double salary = Double.parseDouble(tokens[1]);
    LocalDate hireDate = LocalDate.parse(tokens[2]);
    int year = hireDate.getYear();
    int month = hireDate.getMonthValue();
    int day = hireDate.getDayOfMonth();
    return new Employee(name, salary, year, month, day);
}

```

The parameter of the `split` method is a regular expression describing the separator. We discuss regular expressions in more detail at the end of this chapter. As it happens, the vertical bar character has a special meaning in regular expressions, so it needs to be escaped with a `\` character. That character needs to be escaped by another `\`, yielding the `"\\|"` expression.

The complete program is in [Listing 2.1](#). The static method

[Click here to view code image](#)

```

void writeData(Employee[] e, PrintWriter out)

```

first writes the length of the array, then writes each record. The static method

[Click here to view code image](#)

```

Employee[] readData(BufferedReader in)

```

first reads in the length of the array, then reads in each record. This turns out to be a bit tricky:

[Click here to view code image](#)

```

int n = in.nextInt();
in.nextLine(); // consume newline
Employee[] employees = new Employee[n];
for (int i = 0; i < n; i++)
{
    employees[i] = new Employee();
    employees[i].readData(in);
}

```

The call to `nextInt` reads the array length but not the trailing newline character. We must consume the newline so that the `readData` method can get the next input line when it calls the `nextLine` method.

Listing 2.1 `textFile/TextFileTest.java`

[Click here to view code image](#)

```

1  package textFile;
2
3  import java.io.*;
4  import java.time.*;
5  import java.util.*;
6
7  /**
8   * @version 1.14 2016-07-11
9   * @author Cay Horstmann
10  */
11 public class TextFileTest
12 {
13     public static void main(String[] args) throws IOException
14     {
15         Employee[] staff = new Employee[3];
16
17         staff[0] = new Employee("Carl Cracker", 75000, 1987, 12, 15);
18         staff[1] = new Employee("Harry Hacker", 50000, 1989, 10, 1);
19         staff[2] = new Employee("Tony Tester", 40000, 1990, 3, 15);
20
21         // save all employee records to the file employee.dat
22         try (PrintWriter out = new PrintWriter("employee.dat", "UTF-8"))
23         {
24             writeData(staff, out);
25         }
26
27         // retrieve all records into a new array
28         try (Scanner in = new Scanner(
29             new FileInputStream("employee.dat"), "UTF-8"))
30         {
31             Employee[] newStaff = readData(in);
32
33             // print the newly read employee records
34             for (Employee e : newStaff)
35                 System.out.println(e);
36         }
37     }
38
39     /**
40      * Writes all employees in an array to a print writer
41      * @param employees an array of employees
42      * @param out a print writer
43      */
44     private static void writeData(Employee[] employees, PrintWriter out)
45     throws IOException
46     {
47         // write number of employees
48         out.println(employees.length);
49
50         for (Employee e : employees)
51             writeEmployee(out, e);
52     }
53
54     /**
55      * Reads an array of employees from a scanner
56      * @param in the scanner
57      * @return the array of employees
58      */
59     private static Employee[] readData(Scanner in)
60     {

```

```

60         // retrieve the array size
61         int n = in.nextInt();
62         in.nextLine(); // consume newline
63
64         Employee[] employees = new Employee[n];
65         for (int i = 0; i < n; i++)
66         {
67             employees[i] = readEmployee(in);
68         }
69         return employees;
70     }
71
72     /**
73      * Writes employee data to a print writer
74      * @param out the print writer
75      */
76     public static void writeEmployee(PrintWriter out, Employee e)
77     {
78         out.println(e.getName() + "|" + e.getSalary() + "|" + e.getHireDay());
79     }
80
81     /**
82      * Reads employee data from a buffered reader
83      * @param in the scanner
84      */
85     public static Employee readEmployee(Scanner in)
86     {
87         String line = in.nextLine();
88         String[] tokens = line.split("\\|");
89         String name = tokens[0];
90         double salary = Double.parseDouble(tokens[1]);
91         LocalDate hireDate = LocalDate.parse(tokens[2]);
92         int year = hireDate.getYear();
93         int month = hireDate.getMonthValue();
94         int day = hireDate.getDayOfMonth();
95         return new Employee(name, salary, year, month, day);
96     }
97 }

```

2.2.4 Character Encodings

Input and output streams are for sequences of bytes, but in many cases you will work with texts—that is, sequences of characters. It then matters how characters are encoded into bytes.

Java uses the Unicode standard for characters. Each character or “code point” has a 21-bit integer number. There are different *character encodings*—methods for packaging those 21-bit numbers into bytes.

The most common encoding is UTF-8, which encodes each Unicode code point into a sequence of one to four bytes (see [Table 2.1](#)). UTF-8 has the advantage that the characters of the traditional ASCII character set, which contains all characters used in English, only take up one byte each.

Character Range	Encoding
0...7F	0a ₆ a ₅ a ₄ a ₃ a ₂ a ₁ a ₀
80...7FF	110a ₁₀ a ₉ a ₈ a ₇ a ₆ 10a ₅ a ₄ a ₃ a ₂ a ₁ a ₀
800...FFFF	1110a ₁₅ a ₁₄ a ₁₃ a ₁₂ 10a ₁₁ a ₁₀ a ₉ a ₈ a ₇ a ₆ 10a ₅ a ₄ a ₃ a ₂ a ₁ a ₀
10000...10FFFF	11110a ₂₀ a ₁₉ a ₁₈ 10a ₁₇ a ₁₆ a ₁₅ a ₁₄ a ₁₃ a ₁₂ 10a ₁₁ a ₁₀ a ₉ a ₈ a ₇ a ₆ 10a ₅ a ₄ a ₃ a ₂ a ₁ a ₀

Table 2.1 UTF-8 Encoding

Another common encoding is UTF-16, which encodes each Unicode code point into one or two 16-bit values (see [Table 2.2](#)). This is the encoding used in Java strings. Actually, there are two forms of UTF-16, called “big-endian” and “little-endian.” Consider the 16-bit value 0x2122. In big-endian format, the more significant byte comes first: 0x21 followed by 0x22. In little-endian format, it is the other way around: 0x22 0x21. To indicate which of the two is used, a file can start with the “byte order mark,” the 16-bit quantity 0xFEFF. A reader can use this value to determine the byte order and discard it.

Character Range	Encoding
0...FFFF	a ₁₅ a ₁₄ a ₁₃ a ₁₂ a ₁₁ a ₁₀ a ₉ a ₈ a ₇ a ₆ a ₅ a ₄ a ₃ a ₂ a ₁ a ₀
10000...10FFFF	110110b ₁₉ b ₁₈ b ₁₇ b ₁₆ a ₁₅ a ₁₄ a ₁₃ a ₁₂ a ₁₁ a ₁₀ 110111a ₉ a ₈ a ₇ a ₆ a ₅ a ₄ a ₃ a ₂ a ₁ a ₀ where b ₁₉ b ₁₈ b ₁₇ b ₁₆ = a ₂₀ a ₁₉ a ₁₈ a ₁₇ a ₁₆ - 1

Table 2.2 UTF-16 Encoding



Caution

Some programs, including Microsoft Notepad, add a byte order mark at the beginning of UTF-8 encoded files. Clearly, this is unnecessary since there are no byte ordering issues in UTF-8. But the Unicode standard allows it, and even suggests that it’s a pretty good idea since it leaves little doubt about the encoding. It is supposed to be removed when reading a UTF-8 encoded file. Sadly, Java does not do that, and bug reports against this issue are closed as “will not fix.” Your best bet is to strip out any leading `\uFEFF` that you find in your input.

In addition to the UTF encodings, there are partial encodings that cover a character range suitable for a given user population. For example, ISO 8859-1 is a one-byte code that includes accented characters used in Western European languages. Shift-JIS is a variable-length code for Japanese characters. A large number of these encodings are still in widespread use.

There is no reliable way to automatically detect the character encoding from a stream of bytes. Some API methods let you use the “default charset”—the character encoding preferred

by the operating system of the computer. Is that the same encoding that is used by your source of bytes? These bytes may well originate from a different part of the world. Therefore, you should always explicitly specify the encoding. For example, when reading a web page, check the `Content-Type` header.



Note

The platform encoding is returned by the static method `Charset.defaultCharset`. The static method `Charset.availableCharsets` returns all available `Charset` instances, as a map from canonical names to `Charset` objects.



Caution

The Oracle implementation of Java has a system property `file.encoding` for overriding the platform default. This is not an officially supported property, and it is not consistently followed by all parts of Oracle's implementation of the Java library. You should not set it.

The `StandardCharsets` class has static variables of type `Charset` for the character encodings that every Java virtual machine must support:

[Click here to view code image](#)

```
StandardCharsets.UTF_8  
StandardCharsets.UTF_16  
StandardCharsets.UTF_16BE  
StandardCharsets.UTF_16LE  
StandardCharsets.ISO_8859_1  
StandardCharsets.US_ASCII
```

To obtain the `Charset` for another encoding, use the static `forName` method:

[Click here to view code image](#)

```
Charset shiftJIS = Charset.forName("Shift-JIS");
```

Use the `Charset` object when reading or writing text. For example, you can turn an array of bytes into a string as

[Click here to view code image](#)

```
String str = new String(bytes, StandardCharsets.UTF_8);
```



Tip

Some methods allow you to specify a character encoding with a `Charset` object or a string. Choose the `StandardCharsets` constants, so you don't have to worry about the correct spelling. For example, `new String(bytes, "UTF 8")` is not acceptable and will cause a runtime error.



Caution

Some methods (such as the `String(byte[])` constructor) use the default platform encoding if you don't specify any; others (such as `Files.readAllLines`) use UTF-8.

2.3 Reading and Writing Binary Data

Text format is convenient for testing and debugging because it is humanly readable, but it is not as efficient as transmitting data in binary format. In the following sections, you will learn how to perform input and output with binary data.

2.3.1 The `DataInput` and `DataOutput` interfaces

The `DataOutput` interface defines the following methods for writing a number, a character, a `boolean` value, or a string in binary format:

[Click here to view code image](#)

```
writeChars  
writeByte  
writeInt  
writeShort  
writeLong  
writeFloat  
writeDouble  
writeChar  
writeBoolean  
writeUTF
```

For example, `writeInt` always writes an integer as a 4-byte binary quantity regardless of the number of digits, and `writeDouble` always writes a `double` as an 8-byte binary quantity. The resulting output is not human-readable, but the space needed will be the same for each value of a given type and reading it back in will be faster than parsing text.



Note

There are two different methods of storing integers and floating-point numbers in memory, depending on the processor you are using. Suppose, for example, you are working with a 4-byte `int`, say the decimal number 1234, or 4D2 in hexadecimal ($1234 = 4 \times 256 + 13 \times 16 + 2$). This value can be stored in such a way that the first of the four bytes in memory holds the most significant byte (MSB) of the value: 00 00 04 D2. This is the so-called big-endian method. Or, we can start with the least significant byte (LSB) first: D2 04 00 00. This is called, naturally enough, the little-endian method. For example, the SPARC uses bigendian; the Pentium, little-endian. This can lead to problems. When a file is saved from C or C++ file, the data are saved exactly as the processor stores them. That makes it challenging to move even the simplest data files from one platform to another. In Java, all values are written in the big-endian fashion, regardless of the processor. That makes Java data files platform-independent.

The `writeUTF` method writes string data using a modified version of 8-bit Unicode Transformation Format. Instead of simply using the standard UTF-8 encoding, sequences of Unicode code units are first represented in UTF-16, and then the result is encoded using the UTF-8 rules. This modified encoding is different for characters with codes higher than `0xFFFF`. It is used for backward compatibility with virtual machines that were built when Unicode had not yet grown beyond 16 bits.

Since nobody else uses this modification of UTF-8, you should only use the `writeUTF` method to write strings intended for a Java virtual machine—for example, in a program that generates bytecodes. Use the `writeChars` method for other purposes.

To read the data back in, use the following methods defined in the `DataInput` interface:

[Click here to view code image](#)

```
readInt  
readShort  
readLong  
readFloat  
readDouble  
readChar  
readBoolean  
readUTF
```

The `DataInputStream` class implements the `DataInput` interface. To read binary data from a file, combine a `DataInputStream` with a source of bytes such as a `FileInputStream`:

[Click here to view code image](#)

```
DataInputStream in = new DataInputStream(new FileInputStream("employee.dat"));
```

Similarly, to write binary data, use the `DataOutputStream` class that implements the `DataOutput` interface:

[Click here to view code image](#)

```
DataOutputStream out = new DataOutputStream(new FileOutputStream("employee.dat"));
```

java.io.DataInput 1.0

- `boolean readBoolean()`
- `byte readByte()`
- `char readChar()`
- `double readDouble()`
- `float readFloat()`
- `int readInt()`
- `long readLong()`
- `short readShort()`

reads in a value of the given type.

- `void readFully(byte[] b)`

reads bytes into the array `b`, blocking until all bytes are read.

Parameters: `b` The buffer into which the data are read

- `void readFully(byte[] b, int off, int len)`

reads bytes into the array `b`, blocking until all bytes are read.

Parameters: `b` The buffer into which the data are read

`off` The start offset of the data

`len` The maximum number of bytes to read

- `String readUTF()`

reads a string of characters in the “modified UTF-8” format.

- `int skipBytes(int n)`

skips `n` bytes, blocking until all bytes are skipped.

Parameters: `n` The number of bytes to be skipped

java.io.DataOutput 1.0

- `void writeBoolean(boolean b)`
- `void writeByte(int b)`
- `void writeChar(int c)`
- `void writeDouble(double d)`
- `void writeFloat(float f)`
- `void writeInt(int i)`
- `void writeLong(long l)`
- `void writeShort(int s)`

writes a value of the given type.

- `void writeChars(String s)`

writes all characters in the string.

- `void writeUTF(String s)`

writes a string of characters in the “modified UTF-8” format.

2.3.2 Random-Access Files

The `RandomAccessFile` class lets you read or write data anywhere in a file. Disk files are random-access, but input/output streams that communicate with a network socket are not. You can open a random-access file either for reading only or for both reading and writing; specify the option by using the string “r” (for read access) or “rw” (for read/write access) as the second argument in the constructor.

[Click here to view code image](#)

```
RandomAccessFile in = new RandomAccessFile("employee.dat", "r");
RandomAccessFile inOut = new RandomAccessFile("employee.dat", "rw");
```

When you open an existing file as a `RandomAccessFile`, it does not get deleted.

A random-access file has a *file pointer* that indicates the position of the next byte to be read or written. The `seek` method can be used to set the file pointer to an arbitrary byte position within the file. The argument to `seek` is a `long` integer between zero and the length of the file in bytes. The `getFilePointer` method returns the current position of the file pointer.

The `RandomAccessFile` class implements both the `DataInput` and `DataOutput` interfaces. To read and write from a random-access file, use methods such as `readInt/writeInt` and `readChar/writeChar` that we discussed in the preceding section.

Let’s walk through an example program that stores employee records in a random-access file. Each record will have the same size. This makes it easy to read an arbitrary record. Suppose you want to position the file pointer to the third record. Simply set the file pointer to the appropriate byte position and start reading.

[Click here to view code image](#)

```
long n = 3;
```

```
in.seek((n - 1) * RECORD_SIZE);
Employee e = new Employee();
e.readData(in);
```

If you want to modify the record and save it back into the same location, remember to set the file pointer back to the beginning of the record:

[Click here to view code image](#)

```
in.seek((n - 1) * RECORD_SIZE);
e.writeData(out);
```

To determine the total number of bytes in a file, use the `length` method. The total number of records is the length divided by the size of each record.

[Click here to view code image](#)

```
long nbytes = in.length(); // length in bytes
int nrecords = (int) (nbytes / RECORD_SIZE);
```

Integers and floating-point values have a fixed size in binary format, but we have to work harder for strings. We provide two helper methods to write and read strings of a fixed size. The `writeFixedString` writes the specified number of code units, starting at the beginning of the string. If there are too few code units, the method pads the string, using zero values.

[Click here to view code image](#)

```
public static void writeFixedString(String s, int size, DataOutput out)
    throws IOException
{
    for (int i = 0; i < size; i++)
    {
        char ch = 0;
        if (i < s.length()) ch = s.charAt(i);
        out.writeChar(ch);
    }
}
```

The `readFixedString` method reads characters from the input stream until it has consumed `size` code units or until it encounters a character with a zero value. Then, it skips past the remaining zero values in the input field. For added efficiency, this method uses the `StringBuilder` class to read in a string.

[Click here to view code image](#)

```
public static String readFixedString(int size, DataInput in)
    throws IOException
{
    StringBuilder b = new StringBuilder(size);
    int i = 0;
    boolean more = true;
    while (more && i < size)
    {
        char ch = in.readChar();
        i++;
        if (ch == 0) more = false;
        else b.append(ch);
    }
    in.skipBytes(2 * (size - i));
}
```

```
    return b.toString();  
}
```

We placed the `writeFixedString` and `readFixedString` methods inside the `DataIO` helper class. To write a fixed-size record, we simply write all fields in binary.

[Click here to view code image](#)

```
DataIO.writeFixedString(e.getName(), Employee.NAME_SIZE, out);  
out.writeDouble(e.getSalary());  
LocalDate hireDay = e.getHireDay();  
out.writeInt(hireDay.getYear());  
out.writeInt(hireDay.getMonthValue());  
out.writeInt(hireDay.getDayOfMonth());
```

Reading the data back is just as simple.

[Click here to view code image](#)

```
String name = DataIO.readFixedString(Employee.NAME_SIZE, in);  
double salary = in.readDouble();  
int y = in.readInt();  
int m = in.readInt();  
int d = in.readInt();
```

Let us compute the size of each record. We will use 40 characters for the name strings. Therefore, each record contains 100 bytes:

- 40 characters = 80 bytes for the name
- 1 double = 8 bytes for the salary
- 3 int = 12 bytes for the date

The program shown in [Listing 2.2](#) writes three records into a data file and then reads them from the file in reverse order. To do this efficiently requires random access—we need to get at the last record first.

Listing 2.2 `randomAccess/RandomAccessTest.java`

[Click here to view code image](#)

```
1  package randomAccess;  
2  
3  import java.io.*;  
4  import java.util.*;  
5  import java.time.*;  
6  
7  /**  
8   * @version 1.13 2016-07-11  
9   * @author Cay Horstmann  
10  */  
11  public class RandomAccessTest  
12  {  
13      public static void main(String[] args) throws IOException  
14      {  
15          Employee[] staff = new Employee[3];  
16  
17          staff[0] = new Employee("Carl Cracker", 75000, 1987, 12, 15);  
18          staff[1] = new Employee("Harry Hacker", 50000, 1989, 10, 1);
```

```

19         staff[2] = new Employee("Tony Tester", 40000, 1990, 3, 15);
20
21         try (DataOutputStream out = new DataOutputStream(new
FileOutputStream("employee.dat")))
22         {
23             // save all employee records to the file employee.dat
24             for (Employee e : staff)
25                 writeData(out, e);
26         }
27
28         try (RandomAccessFile in = new RandomAccessFile("employee.dat", "r"))
29         {
30             // retrieve all records into a new array
31
32             // compute the array size
33             int n = (int)(in.length() / Employee.RECORD_SIZE);
34             Employee[] newStaff = new Employee[n];
35
36             // read employees in reverse order
37             for (int i = n - 1; i >= 0; i--)
38             {
39                 newStaff[i] = new Employee();
40                 in.seek(i * Employee.RECORD_SIZE);
41                 newStaff[i] = readData(in);
42             }
43
44             // print the newly read employee records
45             for (Employee e : newStaff)
46                 System.out.println(e);
47         }
48     }
49
50     /**
51     * Writes employee data to a data output
52     * @param out the data output
53     * @param e the employee
54     */
55     public static void writeData(DataOutput out, Employee e) throws
IOException
56     {
57         DataIO.writeFixedString(e.getName(), Employee.NAME_SIZE, out);
58         out.writeDouble(e.getSalary());
59
60         LocalDate hireDay = e.getHireDay();
61         out.writeInt(hireDay.getYear());
62         out.writeInt(hireDay.getMonthValue());
63         out.writeInt(hireDay.getDayOfMonth());
64     }
65
66     /**
67     * Reads employee data from a data input
68     * @param in the data input
69     * @return the employee
70     */
71     public static Employee readData(DataInput in) throws IOException
72     {
73         String name = DataIO.readFixedString(Employee.NAME_SIZE, in);
74         double salary = in.readDouble();
75         int y = in.readInt();
76         int m = in.readInt();

```



```

77         int d = in.readInt();
78         return new Employee(name, salary, y, m - 1, d);
79     }
80 }

```

java.io.RandomAccessFile 1.0

- RandomAccessFile(String file, String mode)

- RandomAccessFile(File file, String mode)

Parameters: file The file to be opened

 mode "r" for read-only mode, "rw" for read/write mode, "rws" for read/write mode with synchronous disk writes of data and metadata for every update, and "rwd" for read/write mode with synchronous disk writes of data only

- long getFilePointer()

returns the current location of the file pointer.

- void seek(long pos)

sets the file pointer to `pos` bytes from the beginning of the file.

- long length()

returns the length of the file in bytes.

2.3.3 ZIP Archives

ZIP archives store one or more files in a (usually) compressed format. Each ZIP archive has a header with information such as the name of each file and the compression method that was used. In Java, you can use a `ZipInputStream` to read a ZIP archive. You need to look at the individual *entries* in the archive. The `getNextEntry` method returns an object of type `ZipEntry` that describes the entry. Pass the entry to the `getInputStream` method of the `ZipInputStream` to obtain an input stream for reading the entry. Then call `closeEntry` to read the next entry. Here is a typical code sequence to read through a ZIP file:

[Click here to view code image](#)

```

ZipInputStream zin = new ZipInputStream(new FileInputStream(zipname));
ZipEntry entry;
while ((entry = zin.getNextEntry()) != null)
{
    InputStream in = zin.getInputStream(entry);
    read the contents of in
    zin.closeEntry();
}
zin.close();

```

To write a ZIP file, use a `ZipOutputStream`. For each entry that you want to place into the ZIP file, create a `ZipEntry` object. Pass the file name to the `ZipEntry` constructor; it sets the other parameters such as file date and decompression method. You can override these settings if you like. Then, call the `putNextEntry` method of the `ZipOutputStream` to begin writing a new

file. Send the file data to the ZIP output stream. When you are done, call `closeEntry`. Repeat for all the files you want to store. Here is a code skeleton:

[Click here to view code image](#)

```
FileOutputStream fout = new FileOutputStream("test.zip");
ZipOutputStream zout = new ZipOutputStream(fout);
for all files
{
    ZipEntry ze = new ZipEntry(filename);
    zout.putNextEntry(ze);
    send data to zout
    zout.closeEntry();
}
zout.close();
```



Note

JAR files (which were discussed in Volume I, Chapter 13) are simply ZIP files with a special entry—the so-called manifest. Use the `JarInputStream` and `JarOutputStream` classes to read and write the manifest entry.

ZIP input streams are a good example of the power of the stream abstraction. When you read data stored in compressed form, you don't need to worry that the data are being decompressed as they are being requested. Moreover, the source of the bytes in a ZIP stream need not be a file—the ZIP data can come from a network connection. In fact, whenever the class loader of an applet reads a JAR file, it reads and decompresses data from the network.



Note

[Section 2.5.8, “ZIP File Systems,” on p. 115](#) shows how to access a ZIP archive without a special API, using the `FileSystem` class of Java SE 7.

`java.util.zip.ZipInputStream` 1.1

- `ZipInputStream(InputStream in)`
creates a `ZipInputStream` that allows you to inflate data from the given `InputStream`.
 - `ZipEntry getNextEntry()`
returns a `ZipEntry` object for the next entry, or `null` if there are no more entries.
 - `void closeEntry()`
closes the current open entry in the ZIP file. You can then read the next entry by using `getNextEntry()`.
-

java.util.zip.ZipOutputStream 1.1

- `ZipOutputStream(OutputStream out)`

creates a `ZipOutputStream` that you can use to write compressed data to the specified `OutputStream`.

- `void putNextEntry(ZipEntry ze)`

writes the information in the given `ZipEntry` to the output stream and positions the stream for the data. The data can then be written by calling the `write()` method.

- `void closeEntry()`

closes the currently open entry in the ZIP file. Use the `putNextEntry` method to start the next entry.

- `void setLevel(int level)`

sets the default compression level of subsequent `DEFLATED` entries. The default value is `Deflater.DEFAULT_COMPRESSION`. Throws an `IllegalArgumentException` if the level is not valid.

Parameters: `level` A compression level, from 0 (`NO_COMPRESSION`) to 9 (`BEST_COMPRESSION`)

- `void setMethod(int method)`

sets the default compression method for this `ZipOutputStream` for any entries that do not specify a method.

Parameters: `method` The compression method, either `DEFLATED` or `STORED`

- ZipEntry(String name)

constructs a zip entry with a given name.

<i>Parameters:</i>	name	The name of the entry
--------------------	------	-----------------------

- long getCrc()

returns the CRC32 checksum value for this `ZipEntry`.

- `String getName()`

returns the name of this entry.

- `long getSize()`

returns the uncompressed size of this entry, or -1 if the uncompressed size is not known.

- `boolean isDirectory()`

returns `true` if this entry is a directory.

- `void setMethod(int method)`

Parameters: method The compression method for the entry; must be either DEFLATED or STORED

- `void setSize(long size)`

sets the size of this entry. Only required if the compression method is `STORED`.

Parameters: size The uncompressed size of this entry

- `void setCrc(long crc)`

sets the CRC32 checksum of this entry. Use the `CRC32` class to compute this checksum. Only required if the compression method is `STORED`.

Parameters: crc The checksum of this entry

java.util.zip.ZipFile 1.1

- `ZipFile(String name)`

- `ZipFile(File file)`

creates a `ZipFile` for reading from the given string or `File` object.

- `Enumeration entries()`

returns an `Enumeration` object that enumerates the `ZipEntry` objects that describe the entries of the `ZipFile`.

- `ZipEntry getEntry(String name)`

returns the entry corresponding to the given name, or `null` if there is no such entry.

Parameters: name The entry name

- `InputStream getInputStream(ZipEntry ze)`

returns an `InputStream` for the given entry.

Parameters: ze A `ZipEntry` in the ZIP file

- `String getName()`

returns the path of this ZIP file.

2.4 Object Input/Output Streams and Serialization

Using a fixed-length record format is a good choice if you need to store data of the same type. However, objects that you create in an object-oriented program are rarely all of the same type. For example, you might have an array called `staff` that is nominally an array of `Employee` records but contains objects that are actually instances of a subclass such as `Manager`. It is certainly possible to come up with a data format that allows you to store such polymorphic collections—but, fortunately, we don't have to. The Java language supports a very general mechanism, called *object serialization*, that makes it possible to write any object to an output stream and read it again later. (You will see in this chapter where the term “serialization” comes from.)

2.4.1 Saving and Loading Serializable Objects

To save object data, you first need to open an `ObjectOutputStream` object:

[Click here to view code image](#)

```
ObjectOutputStream out = new ObjectOutputStream(new
    FileOutputStream("employee.dat"));
```

Now, to save an object, simply use the `writeObject` method of the `ObjectOutputStream` class as in the following fragment:

[Click here to view code image](#)

```
Employee harry = new Employee("Harry Hacker", 50000, 1989, 10, 1);
Manager boss = new Manager("Carl Cracker", 80000, 1987, 12, 15);
out.writeObject(harry);
out.writeObject(boss);
```

To read the objects back in, first get an `ObjectInputStream` object:

[Click here to view code image](#)

```
ObjectInputStream in = new ObjectInputStream(new FileInputStream("employee.dat"));
```

Then, retrieve the objects in the same order in which they were written, using the `readObject` method:

[Click here to view code image](#)

```
Employee e1 = (Employee) in.readObject();  
Employee e2 = (Employee) in.readObject();
```

There is, however, one change you need to make to any class that you want to save to an output stream and restore from an object input stream. The class must implement the `Serializable` interface:

[Click here to view code image](#)

```
class Employee implements Serializable { . . . }
```

The `Serializable` interface has no methods, so you don't need to change your classes in any way. In this regard, it is similar to the `Cloneable` interface that we discussed in Volume I, [Chapter 6](#). However, to make a class cloneable, you still had to override the `clone` method of the `Object` class. To make a class serializable, you do not need to do anything else.



Note

You can write and read only *objects* with the `writeObject/readObject` methods. For primitive type values, use methods such as `writeInt/readInt` or `writeDouble/readDouble`. (The object input/output stream classes implement the `DataInput/DataOutput` interfaces.)

Behind the scenes, an `ObjectOutputStream` looks at all the fields of the objects and saves their contents. For example, when writing an `Employee` object, the name, date, and salary fields are written to the output stream.

However, there is one important situation that we need to consider: What happens when one object is shared by several objects as part of its state?

To illustrate the problem, let us make a slight modification to the `Manager` class. Let's assume that each manager has a secretary:

[Click here to view code image](#)

```
class Manager extends Employee  
{  
    private Employee secretary;  
    . . .  
}
```

Each `Manager` object now contains a reference to the `Employee` object that describes the secretary. Of course, two managers can share the same secretary, as is the case in [Figure 2.5](#) and the following code:

[Click here to view code image](#)

```
harry = new Employee("Harry Hacker", . . .);
Manager carl = new Manager("Carl Cracker", . . .);
carl.setSecretary(harry);
Manager tony = new Manager("Tony Tester", . . .);
tony.setSecretary(harry);
```

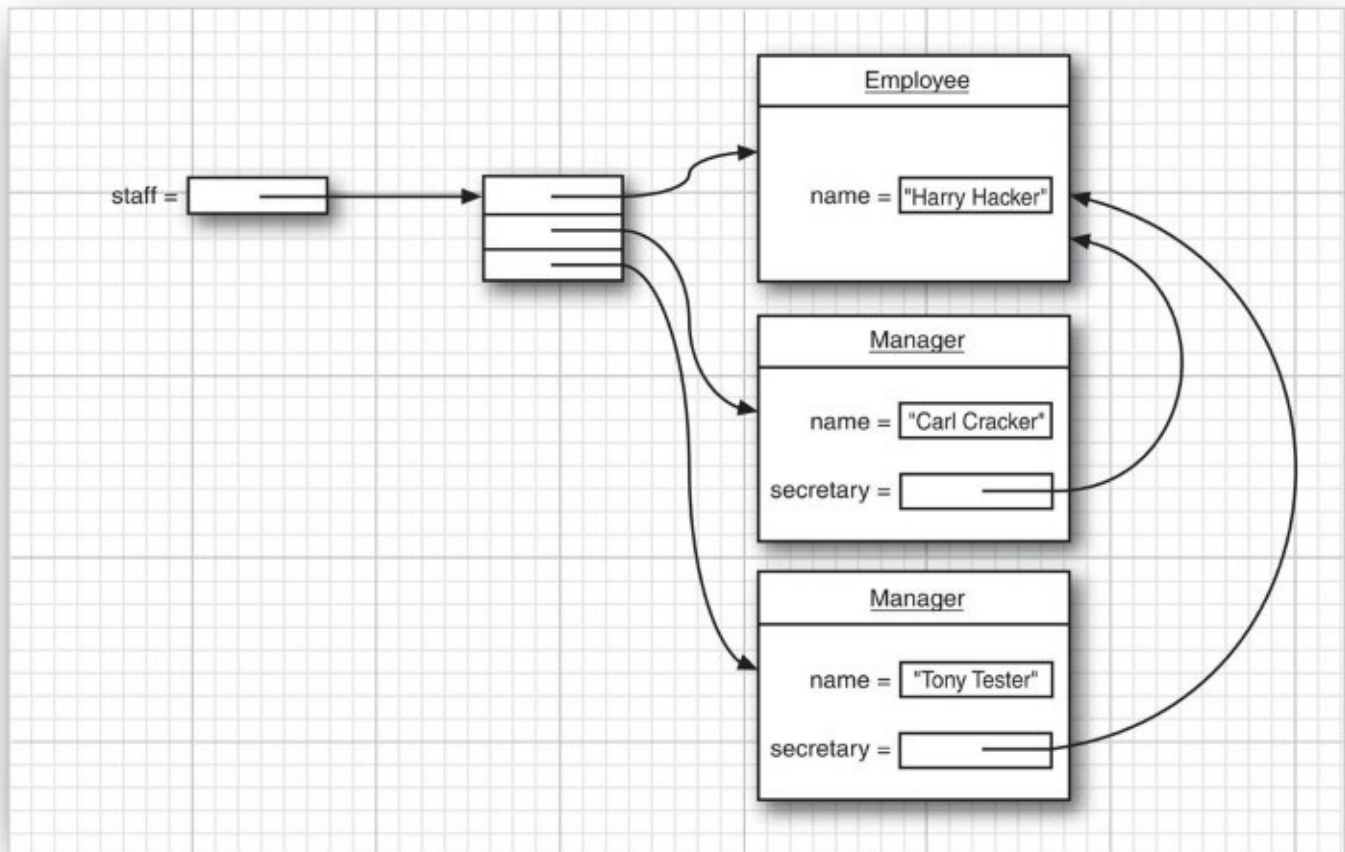


Figure 2.5 Two managers can share a mutual employee.

Saving such a network of objects is a challenge. Of course, we cannot save and restore the memory addresses for the secretary objects. When an object is reloaded, it will likely occupy a completely different memory address than it originally did.

Instead, each object is saved with the *serial number*, hence the name *object serialization* for this mechanism. Here is the algorithm:

1. Associate a serial number with each object reference that you encounter (as shown in [Figure 2.6](#)).

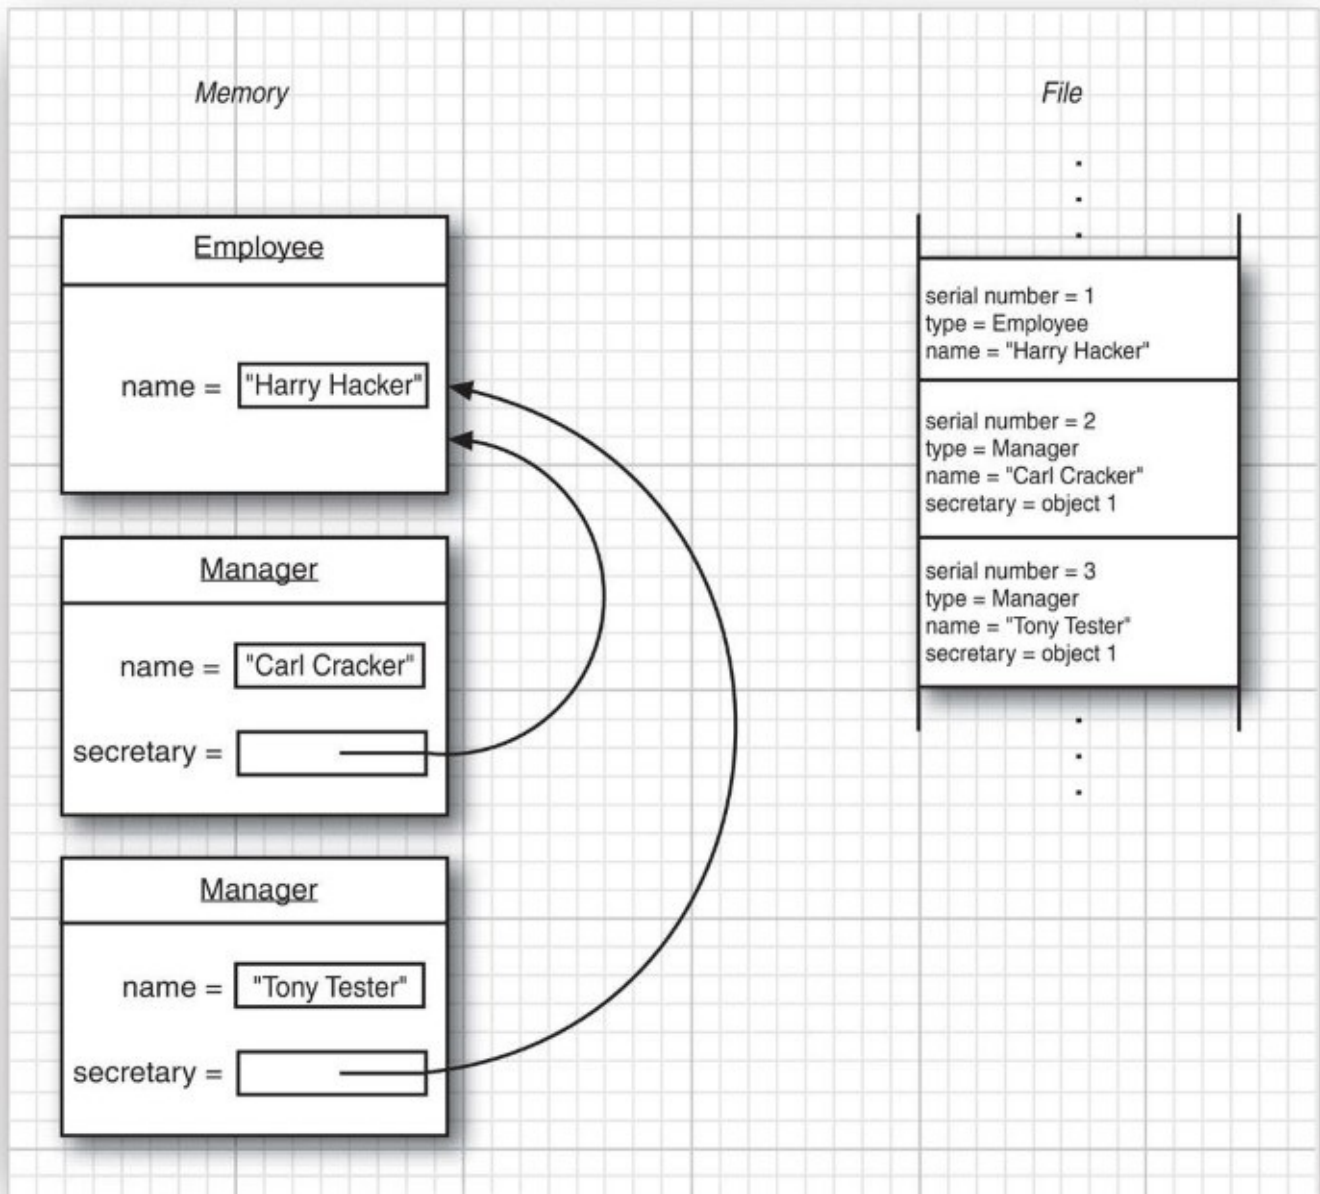


Figure 2.6 An example of object serialization

2. When encountering an object reference for the first time, save the object data to the output stream.
3. If it has been saved previously, just write "same as the previously saved object with serial number x."

When reading back the objects, the procedure is reversed.

1. When an object is specified in an object input stream for the first time, construct it, initialize it with the stream data, and remember the association between the serial number and the object reference.
2. When the tag "same as the previously saved object with serial number x" is encountered, retrieve the object reference for the sequence number.



Note

In this chapter, we will use serialization to save a collection of objects to a disk file and retrieve it exactly as we stored it. Another very important application is the transmittal of a collection of objects across a network connection to another computer. Just as raw memory addresses are meaningless in a file, they are also meaningless when communicating with a different processor. By replacing memory addresses with serial numbers, serialization permits the transport of object collections from one machine to another.

[Listing 2.3](#) is a program that saves and reloads a network of `Employee` and `Manager` objects (some of which share the same employee as a secretary). Note that the secretary object is unique after reloading—when `newStaff[1]` gets a raise, that is reflected in the `secretary` fields of the managers.

Listing 2.3 `objectStream/ObjectStreamTest.java`

[Click here to view code image](#)

```
1  package objectStream;
2
3  import java.io.*;
4
5  /**
6   * @version 1.10 17 Aug 1998
7   * @author Cay Horstmann
8   */
9  class ObjectStreamTest
10 {
11     public static void main(String[] args) throws IOException,
ClassNotFoundException
12     {
13         Employee harry = new Employee("Harry Hacker", 50000, 1989, 10, 1);
14         Manager carl = new Manager("Carl Cracker", 80000, 1987, 12, 15);
15         carl.setSecretary(harry);
16         Manager tony = new Manager("Tony Tester", 40000, 1990, 3, 15);
17         tony.setSecretary(harry);
18
19         Employee[] staff = new Employee[3];
20
21         staff[0] = carl;
22         staff[1] = harry;
23         staff[2] = tony;
24
25         // save all employee records to the file employee.dat
26         try (ObjectOutputStream out = new ObjectOutputStream(new
FileOutputStream("employee.dat")))
27         {
28             out.writeObject(staff);
29         }
30
31         try (ObjectInputStream in = new ObjectInputStream(new
FileInputStream("employee.dat")))
32         {
```

```
33         // retrieve all records into a new array
34
35         Employee[] newStaff = (Employee[]) in.readObject();
36
37         // raise secretary's salary
38         newStaff[1].raiseSalary(10);
39
40         // print the newly read employee records
41         for (Employee e : newStaff)
42             System.out.println(e);
43     }
44 }
45 }
```

java.io.ObjectOutputStream 1.1

- `ObjectOutputStream(OutputStream out)`

creates an `ObjectOutputStream` so that you can write objects to the specified `OutputStream`.

- `void writeObject(Object obj)`

writes the specified object to the `ObjectOutputStream`. This method saves the class of the object, the signature of the class, and the values of any nonstatic, nontransient fields of the class and its superclasses.

java.io.ObjectInputStream 1.1

- `ObjectInputStream(InputStream in)`

creates an `ObjectInputStream` to read back object information from the specified `InputStream`.

- `Object readObject()`

reads an object from the `ObjectInputStream`. In particular, this method reads back the class of the object, the signature of the class, and the values of the nontransient and nonstatic fields of the class and all its superclasses. It does deserializing to allow multiple object references to be recovered.

2.4.2 Understanding the Object Serialization File Format

Object serialization saves object data in a particular file format. Of course, you can use the `writeObject/readObject` methods without having to know the exact sequence of bytes that represents objects in a file. Nonetheless, we found studying the data format extremely helpful for gaining insight into the object serialization process. As the details are somewhat technical, feel free to skip this section if you are not interested in the implementation.

Every file begins with the two-byte “magic number”

AC ED

followed by the version number of the object serialization format, which is currently

00 05

(We use hexadecimal numbers throughout this section to denote bytes.) Then, it contains a sequence of objects, in the order in which they were saved.

String objects are saved as

74 two-byte characters
 length

For example, the string “Harry” is saved as

74 00 05 Harry

The Unicode characters of the string are saved in the “modified UTF-8” format.

When an object is saved, the class of that object must be saved as well. The class description contains

- The name of the class
- The *serial version unique ID*, which is a fingerprint of the data field types and method signatures
- A set of flags describing the serialization method
- A description of the data fields

The fingerprint is obtained by ordering the descriptions of the class, superclass, interfaces, field types, and method signatures in a canonical way, and then applying the so-called Secure Hash Algorithm (SHA) to that data.

SHA is a fast algorithm that gives a “fingerprint” to a larger block of information. This fingerprint is always a 20-byte data packet, regardless of the size of the original data. It is created by a clever sequence of bit operations on the data that makes it essentially 100 percent certain that the fingerprint will change if the information is altered in any way. (For more details on SHA, see, for example, *Cryptography and Network Security, Seventh Edition* by William Stallings, Prentice Hall, 2016.) However, the serialization mechanism uses only the first eight bytes of the SHA code as a class fingerprint. It is still very likely that the class fingerprint will change if the data fields or methods change.

When reading an object, its fingerprint is compared against the current fingerprint of the class. If they don’t match, it means the class definition has changed after the object was written, and an exception is generated. Of course, in practice, classes do evolve, and it might be necessary for a program to read in older versions of objects. We will discuss this in [Section 2.4.5](#), “[Versioning](#),” on [p. 95](#).

Here is how a class identifier is stored:

- 72
- 2-byte length of class name
- Class name
- 8-byte fingerprint
- 1-byte flag
- 2-byte count of data field descriptors

- Data field descriptors
- 78 (end marker)
- Superclass type (70 if none)

The flag byte is composed of three bit masks, defined in `java.io.ObjectStreamConstants`:

[Click here to view code image](#)

```
static final byte SC_WRITE_METHOD = 1;
    // class has a writeObject method that writes additional data
static final byte SC_SERIALIZABLE = 2;
    // class implements the Serializable interface
static final byte SC_EXTERNALIZABLE = 4;
    // class implements the Externalizable interface
```

We discuss the `Externalizable` interface later in this chapter. Externalizable classes supply custom read and write methods that take over the output of their instance fields. The classes that we write implement the `Serializable` interface and will have a flag value of 02. The serializable `java.util.Date` class defines its own `readObject/writeObject` methods and has a flag of 03.

Each data field descriptor has the format:

- 1-byte type code
- 2-byte length of field name
- Field name
- Class name (if the field is an object)

The type code is one of the following:

[Click here to view code image](#)

B	byte
C	char
D	double
F	float
I	int
J	long
L	object
S	short
Z	boolean
[array

When the type code is `L`, the field name is followed by the field type. Class and field name strings do not start with the string code 74, but field types do. Field types use a slightly different encoding of their names—namely, the format used by native methods.

For example, the `salary` field of the `Employee` class is encoded as:

```
D 00 06 salary
```

Here is the complete class descriptor of the `Employee` class:

72 00 08 Employee

E6 D2 86 7D AE AC 18 1B 02

00 03

D 00 06 salary

L 00 07 hireDay

```
74 00 10 Ljava/util/Date;
```

L 00 04 name

```
74 00 12 Ljava/lang/String;
```

78

70

Fingerprint and flags

Number of instance fields

Instance field type and name

Instance field type and name

Instance field class name: Date

Instance field type and name

Instance field class name: String

End marker

No superclass

These descriptors are fairly long. If the *same* class descriptor is needed again in the file, an abbreviated form is used:

71 4-byte serial number

The serial number refers to the previous explicit class descriptor. We discuss the numbering scheme later.

An object is stored as

73	class descriptor	object data
----	------------------	-------------

For example, here is how an `Employee` object is stored:

```
40 E8 6A 00 00 00 00 00 salary field value: double
```

```
73 hireDay field value: new object
```

```
71 00 7E 00 08      Existing class java.util.Date
```

77 08 00 00 00 91 1B 4E B1 80 78 External storage (details later)

```
74 00 0C Harry Hacker          name field value: String
```

As you can see, the data file contains enough information to restore the `Employee` object.

Arrays are saved in the following format:

75	class descriptor	4-byte number of entries	entries
----	------------------	--------------------------	---------

The array class name in the class descriptor is in the same format as that used by native methods (which is slightly different from the format used by class names in other class descriptors). In this format, class names start with an `␣` and end with a semicolon.

For example, an array of three `Employee` objects starts out like this:

75

72 00 0B [LEmployee;

FC BF 36 11 C5 91 11 C7 02

00 00

78

70

00 00 00 03

Array

New class, string length, class name Employee[]

Fingerprint and flags

Number of instance fields

End marker

No superclass

Number of array entries

Note that the fingerprint for an array of `Employee` objects is different from a fingerprint of the `Employee` class itself.

All objects (including arrays and strings) and all class descriptors are given serial numbers as they are saved in the output file. The numbers start at `00 7E 00 00`.

We already saw that a full class descriptor for any given class occurs only once. Subsequent descriptors refer to it. For example, in our previous example, a repeated reference to the `Date` class was coded as

71 00 7E 00 08

The same mechanism is used for objects. If a reference to a previously saved object is written, it is saved in exactly the same way—that is, `71` followed by the serial number. It is always clear from the context whether a particular serial reference denotes a class descriptor or an object.

Finally, a null reference is stored as

70

Here is the commented output of the `ObjectRefTest` program of the preceding section. Run the program, look at a hex dump of its data file `employee.dat`, and compare it with the commented listing. The important lines toward the end of the output show a reference to a previously saved object.

AC ED 00 05	File header
75	Array staff (serial #1)
72 00 0B LEmployee;	New class, string length, class name Employee[] (serial #0)
FC BF 36 11 C5 91 11 C7 02	Fingerprint and flags
00 00	Number of instance fields
78	End marker
70	No superclass
00 00 00 03	Number of array entries
73	staff[0]— new object (serial #7)
72 00 07 Manager	New class, string length, class name (serial #2)
36 06 AE 13 63 8F 59 87 02	Fingerprint and flags
00 01	Number of data fields
L 00 09 secretary	Instance field type and name
74 00 0A LEmployee;	Instance field class name: String (serial #3)
78	End marker
72 00 08 Employee	Superclass: new class, string length, class name (serial #4)
E6 D2 86 7D AE AC 18 18 02	Fingerprint and flags
00 03	Number of instance fields
D 00 06 salary	Instance field type and name
L 00 07 hireDay	Instance field type and name
74 00 10 Ljava/util/Date;	Instance field class name: String (serial #5)
L 00 04 name	Instance field type and name
74 00 12 Ljava/lang/String;	Instance field class name: String (serial #6)
78	End marker
70	No superclass
40 F3 88 00 00 00 00 00	salary field value: double
73	hireDay field value: new object (serial #9)
72 00 0E java.util.Date	New class, string length, class name (serial #8)
68 6A 81 01 4B 59 74 19 03	Fingerprint and flags
00 00	No instance variables
78	End marker
70	No superclass
77 08	External storage, number of bytes
00 00 00 83 E9 39 E0 00	Date
78	End marker
74 00 0C Carl Cracker	name field value: String (serial #10)
73	secretary field value: new object (serial #11)
71 00 7E 00 04	existing class (use serial #4)
40 E8 6A 00 00 00 00 00	salary field value: double
73	hireDay field value: new object (serial #12)
71 00 7E 00 08	Existing class (use serial #8)
77 08	External storage, number of bytes
00 00 00 91 1B 4E 81 80	Date
78	End marker
74 00 0C Harry Hacker	name field value: String (serial #13)
71 00 7E 00 08	staff[1]: existing object (use serial #11)
73	staff[2]: new object (serial #14)
71 00 7E 00 02	Existing class (use serial #2)
40 E3 88 00 00 00 00 00	salary field value: double
73	hireDay field value: new object (serial #15)
71 00 7E 00 08	Existing class (use serial #8)
77 08	External storage, number of bytes
00 00 00 94 6D 3E EC 00 00	Date
78	End marker
74 00 0B Tony Tester	name field value: String (serial #16)
71 00 7E 00 08	secretary field value: existing object (use serial #11)

Of course, studying these codes can be about as exciting as reading a phone book. It is not important to know the exact file format (unless you are trying to create an evil effect by modifying the data), but it is still instructive to know that the serialized format has a detailed description of all the objects that it contains, with sufficient detail to allow reconstruction of both objects and arrays of objects.

What you should remember is this:

- The serialized format contains the types and data fields of all objects.
- Each object is assigned a serial number.
- Repeated occurrences of the same object are stored as references to that serial number.

2.4.3 Modifying the Default Serialization Mechanism

Certain data fields should never be serialized—for example, integer values that store file handles or handles of windows that are only meaningful to native methods. Such information is guaranteed to be useless when you reload an object at a later time or transport it to a different machine. In fact, improper values for such fields can actually cause native methods to crash. Java has an easy mechanism to prevent such fields from ever being serialized: Mark them with the keyword `transient`. You also need to tag fields as `transient` if they belong to nonserializable classes. Transient fields are always skipped when objects are serialized.

The serialization mechanism provides a way for individual classes to add validation or any other desired action to the default read and write behavior. A serializable class can define methods with the signature

[Click here to view code image](#)

```
private void readObject(ObjectInputStream in)
    throws IOException, ClassNotFoundException;
private void writeObject(ObjectOutputStream out)
    throws IOException;
```

Then, the data fields are no longer automatically serialized, and these methods are called instead.

Here is a typical example. A number of classes in the `java.awt.geom` package, such as `Point2D.Double`, are not serializable. Now, suppose you want to serialize a class `LabeledPoint` that stores a `String` and a `Point2D.Double`. First, you need to mark the `Point2D.Double` field as `transient` to avoid a `NotSerializableException`.

[Click here to view code image](#)

```
public class LabeledPoint implements Serializable
{
    private String label;
    private transient Point2D.Double point;
    . . .
}
```

In the `writeObject` method, we first write the object descriptor and the `String` field, `label`, by calling the `defaultWriteObject` method. This is a special method of the `ObjectOutputStream` class that can only be called from within a `writeObject` method of a serializable class. Then we write the point coordinates, using the standard `DataOutput` calls.

[Click here to view code image](#)

```
private void writeObject(ObjectOutputStream out)
    throws IOException
{
    out.defaultWriteObject();
    out.writeDouble(point.getX());
}
```

```
        out.writeDouble(point.getY());
    }
```

In the `readObject` method, we reverse the process:

[Click here to view code image](#)

```
private void readObject(ObjectInputStream in)
    throws IOException
{
    in.defaultReadObject();
    double x = in.readDouble();
    double y = in.readDouble();
    point = new Point2D.Double(x, y);
}
```

Another example is the `java.util.Date` class that supplies its own `readObject` and `writeObject` methods. These methods write the date as a number of milliseconds from the epoch (January 1, 1970, midnight UTC). The `Date` class has a complex internal representation that stores both a `Calendar` object and a millisecond count to optimize lookups. The state of the `Calendar` is redundant and does not have to be saved.

The `readObject` and `writeObject` methods only need to save and load their data fields. They should not concern themselves with superclass data or any other class information.

Instead of letting the serialization mechanism save and restore object data, a class can define its own mechanism. To do this, a class must implement the `Externalizable` interface. This, in turn, requires it to define two methods:

[Click here to view code image](#)

```
public void readExternal(ObjectInputStream in)
    throws IOException, ClassNotFoundException;
public void writeExternal(ObjectOutputStream out)
    throws IOException;
```

Unlike the `readObject` and `writeObject` methods that were described in the previous section, these methods are fully responsible for saving and restoring the entire object, *including the superclass data*. When writing an object, the serialization mechanism merely records the class of the object in the output stream. When reading an externalizable object, the object input stream creates an object with the no-argument constructor and then calls the `readExternal` method. Here is how you can implement these methods for the `Employee` class:

[Click here to view code image](#)

```
public void readExternal(ObjectInput s)
    throws IOException
{
    name = s.readUTF();
    salary = s.readDouble();
    hireDay = LocalDate.ofEpochDay(s.readLong());
}

public void writeExternal(ObjectOutput s)
    throws IOException
{
    s.writeUTF(name);
    s.writeDouble(salary);
}
```

```
s.writeLong(hireDay.toEpochDay());  
}
```



Caution

Unlike the `readObject` and `writeObject` methods, which are private and can only be called by the serialization mechanism, the `readExternal` and `writeExternal` methods are public. In particular, `readExternal` potentially permits modification of the state of an existing object.

2.4.4 Serializing Singletons and Typesafe Enumerations

You have to pay particular attention to serializing and deserializing objects that are assumed to be unique. This commonly happens when you are implementing singletons and typesafe enumerations.

If you use the `enum` construct of the Java language, you need not worry about serialization—it just works. However, suppose you maintain legacy code that contains an enumerated type such as

[Click here to view code image](#)

```
public class Orientation  
{  
    public static final Orientation HORIZONTAL = new Orientation(1);  
    public static final Orientation VERTICAL = new Orientation(2);  
  
    private int value;  
  
    private Orientation(int v) { value = v; }  
}
```

This idiom was common before enumerations were added to the Java language. Note that the constructor is private. Thus, no objects can be created beyond `Orientation.HORIZONTAL` and `Orientation.VERTICAL`. In particular, you can use the `==` operator to test for object equality:

[Click here to view code image](#)

```
if (orientation == Orientation.HORIZONTAL) . . .
```

There is an important twist that you need to remember when a typesafe enumeration implements the `Serializable` interface. The default serialization mechanism is not appropriate. Suppose we write a value of type `Orientation` and read it in again:

[Click here to view code image](#)

```
Orientation original = Orientation.HORIZONTAL;  
ObjectOutputStream out = . . .;  
out.write(original);  
out.close();  
ObjectInputStream in = . . .;  
Orientation saved = (Orientation) in.read();
```

Now the test

[Click here to view code image](#)

```
if (saved == Orientation.HORIZONTAL) . . .
```

will fail. In fact, the `saved` value is a completely new object of the `Orientation` type that is not equal to any of the predefined constants. Even though the constructor is private, the serialization mechanism can create new objects!

To solve this problem, you need to define another special serialization method, called `readResolve`. If the `readResolve` method is defined, it is called after the object is deserialized. It must return an object which then becomes the return value of the `readObject` method. In our case, the `readResolve` method will inspect the `value` field and return the appropriate enumerated constant:

[Click here to view code image](#)

```
protected Object readResolve() throws ObjectStreamException
{
    if (value == 1) return Orientation.HORIZONTAL;
    if (value == 2) return Orientation.VERTICAL;
    throw new ObjectStreamException(); // this shouldn't happen
}
```

Remember to add a `readResolve` method to all typesafe enumerations in your legacy code and to all classes that follow the singleton design pattern.

2.4.5 Versioning

If you use serialization to save objects, you need to consider what happens when your program evolves. Can version 1.1 read the old files? Can the users who still use 1.0 read the files that the new version is producing? Clearly, it would be desirable if object files could cope with the evolution of classes.

At first glance, it seems that this would not be possible. When a class definition changes in any way, its SHA fingerprint also changes, and you know that object input streams will refuse to read in objects with different fingerprints. However, a class can indicate that it is *compatible* with an earlier version of itself. To do this, you must first obtain the fingerprint of the *earlier* version of the class. Use the standalone `serialver` program that is part of the JDK to obtain this number. For example, running

```
serialver Employee
```

prints

[Click here to view code image](#)

```
Employee: static final long serialVersionUID = -1814239825517340645L;
```

If you start the `serialver` program with the `-show` option, the program brings up a graphical dialog box (see [Figure 2.7](#)).

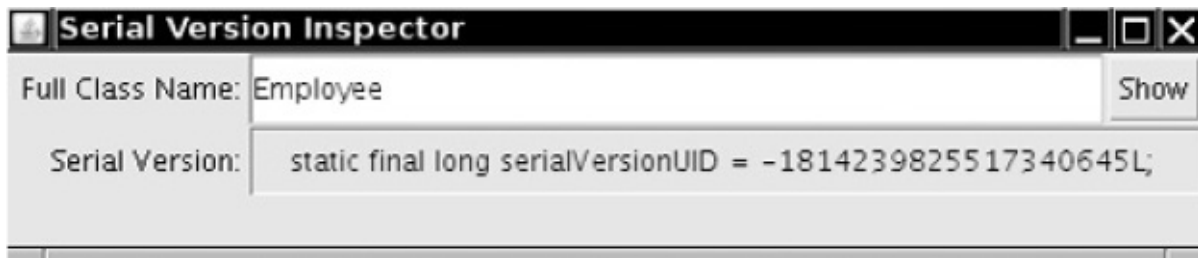


Figure 2.7 The graphical version of the serialver program

All *later* versions of the class must define the `serialVersionUID` constant to the same fingerprint as the original.

[Click here to view code image](#)

```
class Employee implements Serializable // version 1.1
{
    . . .
    public static final long serialVersionUID = -1814239825517340645L;
}
```

When a class has a static data member named `serialVersionUID`, it will not compute the fingerprint manually but will use that value instead.

Once that static data member has been placed inside a class, the serialization system is now willing to read in different versions of objects of that class.

If only the methods of the class change, there is no problem with reading the new object data. However, if the data fields change, you may have problems. For example, the old file object may have more or fewer data fields than the one in the program, or the types of the data fields may be different. In that case, the object input stream makes an effort to convert the serialized object to the current version of the class.

The object input stream compares the data fields of the current version of the class with those of the version in the serialized object. Of course, the object input stream considers only the nontransient and nonstatic data fields. If two fields have matching names but different types, the object input stream makes no effort to convert one type to the other—the objects are incompatible. If the serialized object has data fields that are not present in the current version, the object input stream ignores the additional data. If the current version has data fields that are not present in the serialized object, the added fields are set to their default (`null` for objects, zero for numbers, and `false` for boolean values).

Here is an example. Suppose we have saved a number of employee records on disk, using the original version (1.0) of the class. Now we change the `Employee` class to version 2.0 by adding a data field called `department`. [Figure 2.8](#) shows what happens when a 1.0 object is read into a program that uses 2.0 objects. The `department` field is set to `null`. [Figure 2.9](#) shows the opposite scenario: A program using 1.0 objects reads a 2.0 object. The additional `department` field is ignored.

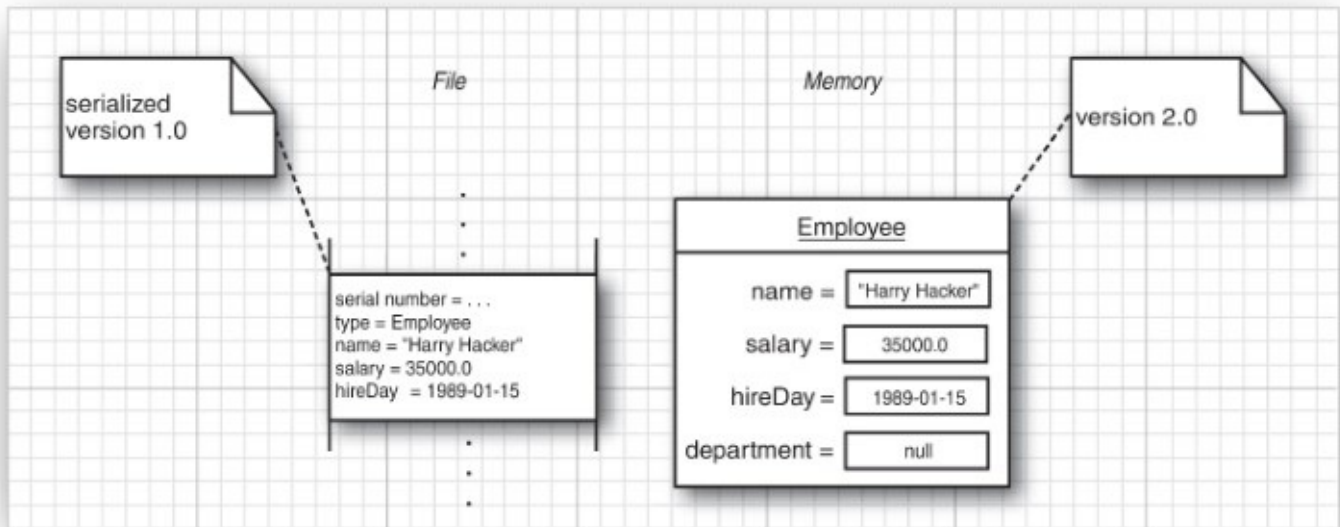


Figure 2.8 Reading an object with fewer data fields

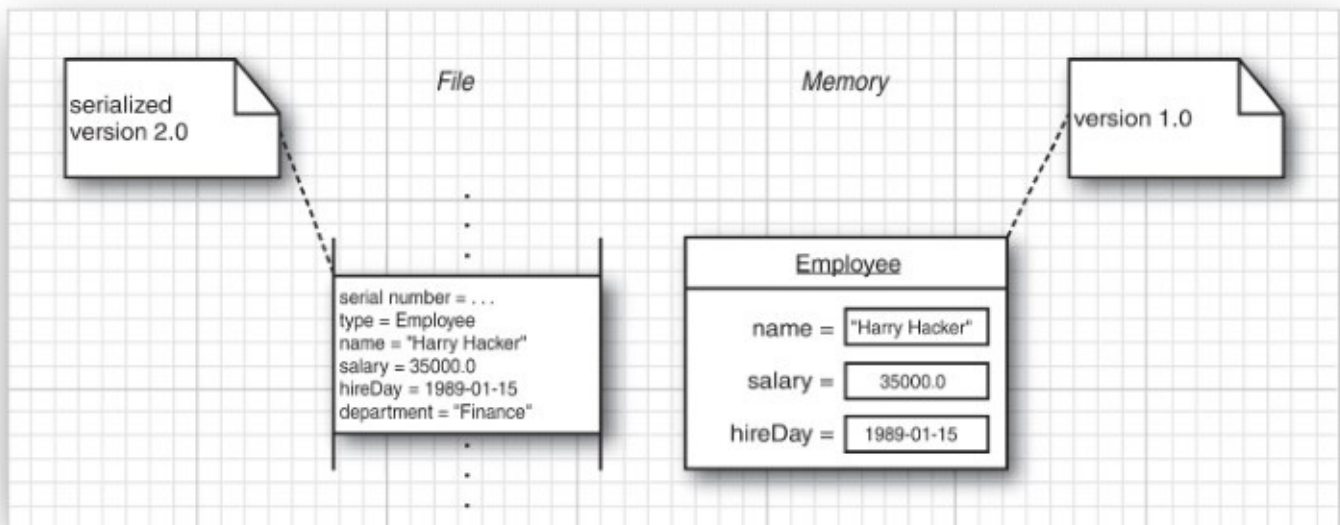


Figure 2.9 Reading an object with more data fields

Is this process safe? It depends. Dropping a data field seems harmless—the recipient still has all the data that it knew how to manipulate. Setting a data field to `null` might not be so safe. Many classes work hard to initialize all data fields in all constructors to non-`null` values, so that the methods don't have to be prepared to handle `null` data. It is up to the class designer to implement additional code in the `readObject` method to fix version incompatibilities or to make sure the methods are robust enough to handle `null` data.

2.4.6 Using Serialization for Cloning

There is an amusing use for the serialization mechanism: It gives you an easy way to clone an object, provided the class is serializable. Simply serialize it to an output stream and then read it back in. The result is a new object that is a deep copy of the existing object. You don't have to write the object to a file—you can use a `ByteArrayOutputStream` to save the data into a byte array.

As [Listing 2.4](#) shows, to get `clone` for free, simply extend the `SerialCloneable` class, and you are done.

You should be aware that this method, although clever, will usually be much slower than a clone method that explicitly constructs a new object and copies or clones the data fields.

Listing 2.4 `serialClone/SerialCloneTest.java`

[Click here to view code image](#)

```
1  package serialClone;
2
3  /**
4   * @version 1.21 13 Jul 2016
5   * @author Cay Horstmann
6   */
7
8  import java.io.*;
9  import java.util.*;
10 import java.time.*;
11
12 public class SerialCloneTest
13 {
14     public static void main(String[] args) throws CloneNotSupportedException
15     {
16         Employee harry = new Employee("Harry Hacker", 35000, 1989, 10, 1);
17         // clone harry
18         Employee harry2 = (Employee) harry.clone();
19
20         // mutate harry
21         harry.raiseSalary(10);
22
23         // now harry and the clone are different
24         System.out.println(harry);
25         System.out.println(harry2);
26     }
27 }
28
29 /**
30  * A class whose clone method uses serialization.
31  */
32 class SerialCloneable implements Cloneable, Serializable
33 {
34     public Object clone() throws CloneNotSupportedException
35     {
36         try {
37             // save the object to a byte array
38             ByteArrayOutputStream bout = new ByteArrayOutputStream();
39             try (ObjectOutputStream out = new ObjectOutputStream(bout))
40             {
```

```

41         out.writeObject(this);
42     }
43
44     // read a clone of the object from the byte array
45     try (InputStream bin = new ByteArrayInputStream(bout.toByteArray()))
46     {
47         ObjectInputStream in = new ObjectInputStream(bin);
48         return in.readObject();
49     }
50 }
51 catch (IOException | ClassNotFoundException e)
52 {
53     CloneNotSupportedException e2 = new CloneNotSupportedException();
54     e2.initCause(e);
55     throw e2;
56 }
57 }
58 }
59
60 /**
61  * The familiar Employee class, redefined to extend the
62  * Serializable class.
63  */
64 class Employee extends Serializable
65 {
66     private String name;
67     private double salary;
68     private LocalDate hireDay;
69
70     public Employee(String n, double s, int year, int month, int day)
71     {
72         name = n;
73         salary = s;
74         hireDay = LocalDate.of(year, month, day);
75     }
76
77     public String getName()
78     {
79         return name;
80     }
81
82     public double getSalary()
83     {
84         return salary;
85     }
86
87     public LocalDate getHireDay()
88     {
89         return hireDay;
90     }
91
92     /**
93      * Raises the salary of this employee.
94      * @byPercent the percentage of the raise
95      */
96     public void raiseSalary(double byPercent)
97     {
98         double raise = salary * byPercent / 100;
99         salary += raise;
100    }

```



```

101
102     public String toString()
103     {
104         return getClass().getName()
105             + "[name=" + name
106             + ",salary=" + salary
107             + ",hireDay=" + hireDay
108             + "]\n";
109     }
110 }

```

2.5 Working with Files

You have learned how to read and write data from a file. However, there is more to file management than reading and writing. The `Path` interface and `Files` class encapsulate the functionality required to work with the file system on the user's machine. For example, the `Files` class can be used to remove or rename the file, or to find out when a file was last modified. In other words, the input/output stream classes are concerned with the contents of files, whereas the classes that we discuss here are concerned with the storage of files on a disk.

The `Path` interface and `Files` class were added in Java SE 7. They are much more convenient to use than the `File` class which dates back all the way to JDK 1.0. We expect them to be very popular with Java programmers and discuss them in depth.

2.5.1 Paths

A `Path` is a sequence of directory names, optionally followed by a file name. The first component of a path may be a *root component* such as `/` or `c:\`. The permissible root components depend on the file system. A path that starts with a root component is *absolute*. Otherwise, it is *relative*. For example, here we construct an absolute and a relative path. For the absolute path, we assume a UNIX-like file system.

[Click here to view code image](#)

```

Path absolute = Paths.get("/home", "harry");
Path relative = Paths.get("myprog", "conf", "user.properties");

```

The static `Paths.get` method receives one or more strings, which it joins with the path separator of the default file system (`/` for a UNIX-like file system, `\` for Windows). It then parses the result, throwing an `InvalidPathException` if the result is not a valid path in the given file system. The result is a `Path` object.

The `get` method can get a single string containing multiple components. For example, you can read a path from a configuration file like this:

[Click here to view code image](#)

```

String baseDir = props.getProperty("base.dir");
// May be a string such as /opt/myprog or c:\Program Files\myprog
Path basePath = Paths.get(baseDir); // OK that baseDir has separators

```



Note

A path does not have to correspond to a file that actually exists. It is merely an abstract sequence of names. As you will see in the next section, when you want to create a file, you first make a path and then call a method to create the corresponding file.

It is very common to combine or *resolve* paths. The call `p.resolve(q)` returns a path according to these rules:

- If `q` is absolute, then the result is `q`.
- Otherwise, the result is “`p` then `q`,” according to the rules of the file system.

For example, suppose your application needs to find its working directory relative to a given base directory that is read from a configuration file, as in the preceding example.

[Click here to view code image](#)

```
Path workRelative = Paths.get("work");
Path workPath = basePath.resolve(workRelative);
```

There is a shortcut for the `resolve` method that takes a string instead of a path:

[Click here to view code image](#)

```
Path workPath = basePath.resolve("work");
```

There is a convenience method `resolveSibling` that resolves against a path's parent, yielding a sibling path. For example, if `workPath` is `/opt/myapp/work`, the call

[Click here to view code image](#)

```
Path tempPath = workPath.resolveSibling("temp");
```

creates `/opt/myXapp/temp`.

The opposite of `resolve` is `relativize`. The call `p.relativize(r)` yields the path `q` which, when resolved with `p`, yields `r`. For example, relativizing `"/home/harry"` against `"/home/fred/input.txt"` yields `"../fred/input.txt"`. Here, we assume that `..` denotes the parent directory in the file system.

The `normalize` method removes any redundant `.` and `..` components (or whatever the file system may deem redundant). For example, normalizing the path `/home/harry/ ../fred/ ./input.txt` yields `/home/fred/input.txt`.

The `toAbsolutePath` method yields the absolute path of a given path, starting at a root component, such as `/home/fred/input.txt` or `c:\Users\fred\input.txt`.

The `Path` interface has many useful methods for taking paths apart. This code sample shows some of the most useful ones:

[Click here to view code image](#)

```
Path p = Paths.get("/home", "fred", "myprog.properties");
Path parent = p.getParent(); // the path /home/fred
Path file = p.getFileName(); // the path myprog.properties
Path root = p.getRoot(); // the path /
```

As you have already seen in Volume I, you can construct a `Scanner` from a `Path` object:

[Click here to view code image](#)

```
Scanner in = new Scanner(Paths.get("/home/fred/input.txt"));
```



Note

Occasionally, you may need to inter-operate with legacy APIs that use the `File` class instead of the `Path` interface. The `Path` interface has a `toFile` method, and the `File` class has a `toPath` method.

`java.nio.file.Paths` 7

- `static Path get(String first, String... more)`
makes a path by joining the given strings.
-
-

`java.nio.file.Path` 7

- `Path resolve(Path other)`
 - `Path resolve(String other)`
if `other` is absolute, returns `other`; otherwise, returns the path obtained by joining this and `other`.
 - `Path resolveSibling(Path other)`
 - `Path resolveSibling(String other)`
if `other` is absolute, returns `other`; otherwise, returns the path obtained by joining the parent of this and `other`.
 - `Path relativize(Path other)`
returns the relative path that, when resolved with this, yields `other`.
 - `Path normalize()`
removes redundant path elements such as `.` and `..`
 - `Path toAbsolutePath()`
returns an absolute path that is equivalent to this path.
 - `Path getParent()`
returns the parent, or `null` if this path has no parent.
 - `Path getFileName()`
returns the last component of this path, or `null` if this path has no components.
 - `Path getRoot()`
returns the root component of this path, or `null` if this path has no root components.
 - `toFile()`
makes a `File` from this path.
-
-

java.io.File 1.0

- `Path toPath()` ⁷

makes a `Path` from this file.

2.5.2 Reading and Writing Files

The `Files` class makes quick work of common file operations. For example, you can easily read the entire contents of a file:

[Click here to view code image](#)

```
byte[] bytes = Files.readAllBytes(path);
```

If you want to read the file as a string, call `readAllBytes` followed by

[Click here to view code image](#)

```
String content = new String(bytes, charset);
```

But if you want the file as a sequence of lines, call

[Click here to view code image](#)

```
List<String> lines = Files.readAllLines(path, charset);
```

Conversely, if you want to write a string, call

[Click here to view code image](#)

```
Files.write(path, content.getBytes(charset));
```

To append to a given file, use

[Click here to view code image](#)

```
Files.write(path, content.getBytes(charset), StandardOpenOption.APPEND);
```

You can also write a collection of lines with

```
Files.write(path, lines);
```

These simple methods are intended for dealing with text files of moderate length. If your files are large or binary, you can still use the familiar input/output streams or readers/writers:

[Click here to view code image](#)

```
InputStream in = Files.newInputStream(path);
OutputStream out = Files.newOutputStream(path);
Reader in = Files.newBufferedReader(path, charset);
Writer out = Files.newBufferedWriter(path, charset);
```

These convenience methods save you from dealing with `FileInputStream`, `FileOutputStream`, `BufferedReader`, or `BufferedWriter`.

java.nio.file.Files 7

- static byte[] readAllBytes(Path path)
 - static List<String> readAllLines(Path path, Charset charset)
reads the contents of a file.
 - static Path write(Path path, byte[] contents, OpenOption... options)
 - static Path write(Path path, Iterable<? extends CharSequence> contents, OpenOption options)
writes the given contents to a file and returns path.
 - static InputStream newInputStream(Path path, OpenOption... options)
 - static OutputStream newOutputStream(Path path, OpenOption... options)
 - static BufferedReader newBufferedReader(Path path, Charset charset)
 - static BufferedWriter newBufferedWriter(Path path, Charset charset, OpenOption... options)
opens a file for reading or writing.
-

2.5.3 Creating Files and Directories

To create a new directory, call

```
Files.createDirectory(path);
```

All but the last component in the path must already exist. To create intermediate directories as well, use

```
Files.createDirectories(path);
```

You can create an empty file with

```
Files.createFile(path);
```

The call throws an exception if the file already exists. The check for existence and creation are atomic. If the file doesn't exist, it is created before anyone else has a chance to do the same.

There are convenience methods for creating a temporary file or directory in a given or system-specific location.

[Click here to view code image](#)

```
Path newPath = Files.createTempFile(dir, prefix, suffix);  
Path newPath = Files.createTempFile(prefix, suffix);  
Path newPath = Files.createTempDirectory(dir, prefix);  
Path newPath = Files.createTempDirectory(prefix);
```

Here, `dir` is a `Path`, and `prefix/suffix` are strings which may be `null`. For example, the call `Files.createTempFile(null, ".txt")` might return a path such as `/tmp/1234405522364837194.txt`.

When you create a file or directory, you can specify attributes, such as owners or permissions. However, the details depend on the file system, and we won't cover them here.

java.nio.file.Files 7

- static Path `createFile(Path path, FileAttribute<?>... attrs)`
- static Path `createDirectory(Path path, FileAttribute<?>... attrs)`
- static Path `createDirectories(Path path, FileAttribute<?>... attrs)`

creates a file or directory. The `createDirectories` method creates any intermediate directories as well.

- static Path `createTempFile(String prefix, String suffix, FileAttribute<?>... attrs)`
- static Path `createTempFile(Path parentDir, String prefix, String suffix, FileAttribute<?>... attrs)`
- static Path `createTempDirectory(String prefix, FileAttribute<?>... attrs)`
- static Path `createTempDirectory(Path parentDir, String prefix, FileAttribute<?>... attrs)`

creates a temporary file or directory, in a location suitable for temporary files or in the given parent directory. Returns the path to the created file or directory.

2.5.4 Copying, Moving, and Deleting Files

To copy a file from one location to another, simply call

```
Files.copy(fromPath, toPath);
```

To move the file (that is, copy and delete the original), call

```
Files.move(fromPath, toPath);
```

The copy or move will fail if the target exists. If you want to overwrite an existing target, use the `REPLACE_EXISTING` option. If you want to copy all file attributes, use the `COPY_ATTRIBUTES` option. You can supply both like this:

[Click here to view code image](#)

```
Files.copy(fromPath, toPath, StandardCopyOption.REPLACE_EXISTING,  
StandardCopyOption.COPY_ATTRIBUTES);
```

You can specify that a move should be atomic. Then you are assured that either the move completed successfully, or the source continues to be present. Use the `ATOMIC_MOVE` option:

[Click here to view code image](#)

```
Files.move(fromPath, toPath, StandardCopyOption.ATOMIC_MOVE);
```

You can also copy an input stream to a `Path`, which just means saving the input stream to disk. Similarly, you can copy a `Path` to an output stream. Use the following calls:

[Click here to view code image](#)

```
Files.copy(inputStream, toPath);  
Files.copy(fromPath, outputStream);
```

As with the other calls to `copy`, you can supply copy options as needed.

Finally, to delete a file, simply call

```
Files.delete(path);
```

This method throws an exception if the file doesn’t exist, so instead you may want to use

[Click here to view code image](#)

```
boolean deleted = Files.deleteIfExists(path);
```

The deletion methods can also be used to remove an empty directory.

See [Table 2.3](#) for a summary of the options that are available for file operations.

Option	Description
StandardOpenOption; use with newBufferedWriter, newInputStream, newOutputStream, write	
READ	Open for reading
WRITE	Open for writing
APPEND	If opened for writing, append to the end of the file
TRUNCATE_EXISTING	If opened for writing, remove existing contents
CREATE_NEW	Create a new file and fail if it exists
CREATE	Atomically create a new file if it doesn’t exist
DELETE_ON_CLOSE	Make a “best effort” to delete the file when it is closed
SPARSE	A hint to the file system that this file will be sparse
DSYNC SYNC	Requires that each update to the file data data and metadata be written synchronously to the storage device
StandardCopyOption; use with copy, move	
ATOMIC_MOVE	Move the file atomically
COPY_ATTRIBUTES	Copy the file attributes
REPLACE_EXISTING	Replace the target if it exists
LinkOption; use with all of the above methods and exists, isDirectory, isRegularFile	
NOFOLLOW_LINKS	Do not follow symbolic links
FileVisitOption; use with find, walk, walkFileTree	
FOLLOW_LINKS	Follow symbolic links

Table 2.3 Standard Options for File Operations

java.nio.file.Files 7

- static Path copy(Path from, Path to, CopyOption... options)
 - static Path move(Path from, Path to, CopyOption... options)
copies or moves from to the given target location and returns to.
 - static long copy(InputStream from, Path to, CopyOption... options)
 - static long copy(Path from, OutputStream to, CopyOption... options)
copies from an input stream to a file, or from a file to an output stream, returning the number of bytes copied.
 - static void delete(Path path)
 - static boolean deleteIfExists(Path path)
deletes the given file or empty directory. The first method throws an exception if the file or directory doesn't exist. The second method returns false in that case.
-

2.5.5 Getting File Information

The following static methods return a `boolean` value to check a property of a path:

- `exists`
- `isHidden`
- `isReadable`, `isWritable`, `isExecutable`
- `isRegularFile`, `isDirectory`, `isSymbolicLink`

The `size` method returns the number of bytes in a file.

[Click here to view code image](#)

```
long fileSize = Files.size(path);
```

The `getOwner` method returns the owner of the file, as an instance of `java.nio.file.attribute.UserPrincipal`.

All file systems report a set of basic attributes, encapsulated by the `BasicFileAttributes` interface, which partially overlaps with that information. The basic file attributes are

- The times at which the file was created, last accessed, and last modified, as instances of the class `java.nio.file.attribute.FileTime`
- Whether the file is a regular file, a directory, a symbolic link, or none of these
- The file size
- The file key—an object of some class, specific to the file system, that may or may not uniquely identify a file

To get these attributes, call

[Click here to view code image](#)

```
BasicFileAttributes attributes = Files.readAttributes(path,  
BasicFileAttributes.class);
```

If you know that the user's file system is POSIX-compliant, you can instead get an instance of

PosixFileAttributes:

[Click here to view code image](#)

```
PosixFileAttributes attributes = Files.readAttributes(path,  
PosixFileAttributes.class);
```

Then you can find out the group owner and the owner, group, and world access permissions of the file. We won't dwell on the details since so much of this information is not portable across operating systems.

***java.nio.file.Files* 7**

- static boolean exists(Path path)
- static boolean isHidden(Path path)
- static boolean isReadable(Path path)
- static boolean isWritable(Path path)
- static boolean isExecutable(Path path)
- static boolean isRegularFile(Path path)
- static boolean isDirectory(Path path)
- static boolean isSymbolicLink(Path path)

checks for the given property of the file given by the path.

- static long size(Path path)
gets the size of the file in bytes.
 - A readAttributes(Path path, Class<A> type, LinkOption... options)
reads the file attributes of type A.
-

***java.nio.file.attribute.BasicFileAttributes* 7**

- FileTime creationTime()
- FileTime lastAccessTime()
- FileTime lastModifiedTime()
- boolean isRegularFile()
- boolean isDirectory()
- boolean isSymbolicLink()
- long size()
- Object fileKey()

gets the requested attribute.

2.5.6 Visiting Directory Entries

The static `Files.list` method returns a `Stream<Path>` that reads the entries of a directory. The directory is read lazily, making it possible to efficiently process directories with huge numbers of entries.

Since reading a directory involves a system resource that needs to be closed, you should use a `try` block:

[Click here to view code image](#)

```
try (Stream<Path> entries = Files.list(pathToDirectory))
{
    . . .
}
```

The `list` method does not enter subdirectories. To process all descendants of a directory, use the `Files.walk` method instead.

[Click here to view code image](#)

```
try (Stream<Path> entries = Files.walk(pathToRoot))
{
    // Contains all descendants, visited in depth-first order
}
```

Here is a sample traversal of the unzipped `src.zip` tree:

[Click here to view code image](#)

```
java
java/nio
java/nio/DirectCharBufferU.java
java/nio/ByteBufferAsShortBufferRL.java
java/nio/MappedByteBuffer.java
. . .
java/nio/ByteBufferAsDoubleBufferB.java
java/nio/charset
java/nio/charset/CoderMalfunctionError.java
java/nio/charset/CharsetDecoder.java
java/nio/charset/UnsupportedCharsetException.java
java/nio/charset/spi
java/nio/charset/spi/CharsetProvider.java
java/nio/charset/StandardCharsets.java
java/nio/charset/Charset.java
. . .
java/nio/charset/CoderResult.java
java/nio/HeapFloatBufferR.java
. . .
```

As you can see, whenever the traversal yields a directory, it is entered before continuing with its siblings.

You can limit the depth of the tree that you want to visit by calling `Files.walk(pathToRoot, depth)`. Both `walk` methods have a varargs parameter of type `FileVisitOption...`, but there is only one option you can supply: `FOLLOW_LINKS` to follow symbolic links.



Note

If you filter the paths returned by `walk` and your filter criterion involves the file attributes stored with a directory, such as size, creation time, or type (file, directory, symbolic link), then use the `find` method instead of `walk`. Call that method with a predicate function that accepts a path and a `BasicFileAttributes` object. The only advantage is efficiency. Since the directory is being read anyway, the attributes are readily available.

This code fragment uses the `Files.walk` method to copy one directory to another:

[Click here to view code image](#)

```
Files.walk(source).forEach(p ->
{
    try
    {
        Path q = target.resolve(source.relativize(p));
        if (Files.isDirectory(p))
            Files.createDirectory(q);
        else
            Files.copy(p, q);
    }
    catch (IOException ex)
    {
        throw new UncheckedIOException(ex);
    }
});
```

Unfortunately, you cannot easily use the `Files.walk` method to delete a tree of directories since you need to first delete the children before deleting the parent. The next section shows you how to overcome that problem.

2.5.7 Using Directory Streams

As you saw in the preceding section, the `Files.walk` method produces a `Stream<Path>` that traverses the descendants of a directory. Sometimes, you need more fine-grained control over the traversal process. In that case, use the `Files.newDirectoryStream` object instead. It yields a `DirectoryStream`. Note that this is not a subinterface of `java.util.stream.Stream` but an interface that is specialized for directory traversal. It is a subinterface of `Iterable` so that you can use directory stream in an enhanced for loop. Here is the usage pattern:

[Click here to view code image](#)

```
try (DirectoryStream<Path> entries = Files.newDirectoryStream(dir))
{
    for (Path entry : entries)
        Process entries
}
```

The try-with-resources block ensures that the directory stream is properly closed.

There is no specific order in which the directory entries are visited.

You can filter the files with a glob pattern:

[Click here to view code image](#)

```
try (DirectoryStream<Path> entries = Files.newDirectoryStream(dir, "*.java"))
```

[Table 2.4](#) shows all glob patterns.

Pattern	Description	Example
*	Matches zero or more characters of a path component	*.java matches all Java files in the current directory
**	Matches zero or more characters, crossing directory boundaries	**.java matches all Java files in any subdirectory
?	Matches one character	????.java matches all four-character Java files (not counting the extension)
[. . .]	Matches a set of characters. You can use hyphens [0-9] and negation [!0-9].	Test[0-9A-F].java matches Testx.java, where x is one hexadecimal digit
{. . .}	Matches alternatives, separated by commas	*.{java,class} matches all Java and class files
\	Escapes any of the above as well as \	*** matches all files with a * in their name

Table 2.4 Glob Patterns



Caution

If you use the glob syntax on Windows, you have to escape backslashes *twice*: once for the glob syntax, and once for the Java string syntax: `Files.newDirectoryStream(dir, "C:\\\\")`.

If you want to visit all descendants of a directory, call the `walkFileTree` method instead and supply an object of type `FileVisitor`. That object gets notified

- **When a file is encountered:** `FileVisitResult visitFile(T path, BasicFileAttributes attrs)`
- **Before a directory is processed:** `FileVisitResult preVisitDirectory(T dir, IOException ex)`
- **After a directory is processed:** `FileVisitResult postVisitDirectory(T dir, IOException ex)`
- **When an error occurred trying to visit a file or directory, such as trying to open a directory without the necessary permissions:** `FileVisitResult visitFileFailed(T path, IOException ex)`

In each case, you can specify whether you want to

- Continue visiting the next file: `FileVisitResult.CONTINUE`
- Continue the walk, but without visiting the entries in this directory:
`FileVisitResult.SKIP_SUBTREE`
- Continue the walk, but without visiting the siblings of this file:
`FileVisitResult.SKIP_SIBLINGS`
- Terminate the walk: `FileVisitResult.TERMINATE`

If any of the methods throws an exception, the walk is also terminated, and that exception is thrown from the `walkFileTree` method.



Note

The `FileVisitor` interface is a generic type, but it isn't likely that you'll ever want something other than a `FileVisitor<Path>`. The `walkFileTree` method is willing to accept a `FileVisitor<? super Path>`, but `Path` does not have an abundance of supertypes.

A convenience class `SimpleFileVisitor` implements the `FileVisitor` interface. All but the `visitFileFailed` method do nothing and continue. The `visitFileFailed` method throws the exception that caused the failure, thereby terminating the visit.

For example, here is how you can print out all subdirectories of a given directory:

[Click here to view code image](#)

```
Files.walkFileTree(Paths.get("/"), new SimpleFileVisitor<Path>()
{
    public FileVisitResult preVisitDirectory(Path path, BasicFileAttributes
attrs)
        throws IOException
    {
        System.out.println(path);
        return FileVisitResult.CONTINUE;
    }

    public FileVisitResult postVisitDirectory(Path dir, IOException exc)
    {
        return FileVisitResult.CONTINUE;
    }

    public FileVisitResult visitFileFailed(Path path, IOException exc) throws
IOException
    {
        return FileVisitResult.SKIP_SUBTREE;
    }
});
```

Note that we need to override `postVisitDirectory` and `visitFileFailed`. Otherwise, the visit would fail as soon as it encounters a directory that it's not allowed to open or a file that it's not allowed to access.

Also note that the attributes of the path are passed as a parameter to the `preVisitDirectory` and `visitFile` method. The visitor already had to make an OS call to get the attributes, since it

needs to distinguish between files and directories. This way, you don't need to make another call.

The other methods of the `FileVisitor` interface are useful if you need to do some work when entering or leaving a directory. For example, when you delete a directory tree, you need to remove the current directory after you have removed all of its files. Here is the complete code for deleting a directory tree:

[Click here to view code image](#)

```
// Delete the directory tree starting at root
Files.walkFileTree(root, new SimpleFileVisitor<Path>()
{
    public FileVisitResult visitFile(Path file, BasicFileAttributes attrs)
throws IOException
    {
        Files.delete(file);
        return FileVisitResult.CONTINUE;
    }

    public FileVisitResult postVisitDirectory(Path dir, IOException e) throws
IOException
    {
        if (e != null) throw e;
        Files.delete(dir);
        return FileVisitResult.CONTINUE;
    }
});
```

java.nio.file.Files 7

- `static DirectoryStream<Path> newDirectoryStream(Path path)`
 - `static DirectoryStream<Path> newDirectoryStream(Path path, String glob)`
gets an iterator over the files and directories in a given directory. The second method only accepts those entries matching the given glob pattern.
 - `static Path walkFileTree(Path start, FileVisitor<? super Path> visitor)`
walks all descendants of the given path, applying the visitor to all descendants.
-

`java.nio.file.SimpleFileVisitor<T>` 7

- `static FileVisitResult visitFile(T path, BasicFileAttributes attrs)`
is called when a file or directory is visited, returns one of `CONTINUE`, `SKIP_SUBTREE`, `SKIP_SIBLINGS`, or `TERMINATE`. The default implementation does nothing and continues.
 - `static FileVisitResult preVisitDirectory(T dir, BasicFileAttributes attrs)`
 - `static FileVisitResult postVisitDirectory(T dir, BasicFileAttributes attrs)`
are called before and after visiting a directory. The default implementation does nothing and continues.
 - `static FileVisitResult visitFileFailed(T path, IOException exc)`
is called if an exception was thrown in an attempt to get information about the given file. The default implementation rethrows the exception, which causes the visit to terminate with that exception. Override the method if you want to continue.
-

2.5.8 ZIP File Systems

The `Paths` class looks up paths in the default file system—the files on the user’s local disk. You can have other file systems. One of the more useful ones is a *ZIP file system*. If `zipname` is the name of a ZIP file, then the call

[Click here to view code image](#)

```
FileSystem fs = FileSystems.newFileSystem(Paths.get(zipname), null);
```

establishes a file system that contains all files in the ZIP archive. It’s an easy matter to copy a file out of that archive if you know its name:

[Click here to view code image](#)

```
Files.copy(fs.getPath(sourceName), targetPath);
```

Here, `fs.getPath` is the analog of `Paths.get` for an arbitrary file system.

To list all files in a ZIP archive, walk the file tree:

[Click here to view code image](#)

```
FileSystem fs = FileSystems.newFileSystem(Paths.get(zipname), null);
Files.walkFileTree(fs.getPath("/"), new SimpleFileVisitor<Path>()
{
    public FileVisitResult visitFile(Path file, BasicFileAttributes attrs)
    throws IOException
    {
        System.out.println(file);
        return FileVisitResult.CONTINUE;
    }
});
```

That is nicer than the API described in [Section 2.3.3](#), “[ZIP Archives](#),” on [p. 77](#) which required a set of new classes just to deal with ZIP archives.

`java.nio.file.FileSystem` 7

- `static FileSystem newFileSystem(Path path, ClassLoader loader)`
iterates over the installed file system providers and, provided that `loader` is not `null`, the file systems that the given class loader can load. Returns the file system created by the first file system provider that accepts the given path. By default, there is a provider for ZIP file systems that accepts files whose names end in `.zip` or `.jar`.
-

`java.nio.file.FileSystem` 7

- `static Path getPath(String first, String... more)`
makes a path by joining the given strings.
-

2.6 Memory-Mapped Files

Most operating systems can take advantage of a virtual memory implementation to “map” a file, or a region of a file, into memory. Then the file can be accessed as if it were an in-memory array, which is much faster than the traditional file operations.

2.6.1 Memory-Mapped File Performance

At the end of this section, you can find a program that computes the CRC32 checksum of a file using traditional file input and a memory-mapped file. On one machine, we got the timing data shown in [Table 2.5](#) when computing the checksum of the 37MB file `rt.jar` in the `jre/lib` directory of the JDK.

Method	Time
Plain input stream	110 seconds
Buffered input stream	9.9 seconds
Random access file	162 seconds
Memory-mapped file	7.2 seconds

Table 2.5 Timing Data for File Operations

As you can see, on this particular machine, memory mapping is a bit faster than using buffered sequential input and dramatically faster than using a `RandomAccessFile`.

Of course, the exact values will differ greatly from one machine to another, but it is obvious that the performance gain, compared to random access, can be substantial. For sequential reading of files of moderate size, on the other hand, there is no reason to use memory mapping.

The `java.nio` package makes memory mapping quite simple. Here is what you do.

First, get a *channel* for the file. A channel is an abstraction for a disk file that lets you access

operating system features such as memory mapping, file locking, and fast data transfers between files.

[Click here to view code image](#)

```
FileChannel channel = FileChannel.open(path, options);
```

Then, get a `ByteBuffer` from the channel by calling the `map` method of the `FileChannel` class. Specify the area of the file that you want to map and a *mapping mode*. Three modes are supported:

- `FileChannel.MapMode.READ_ONLY`: The resulting buffer is read-only. Any attempt to write to the buffer results in a `ReadOnlyBufferException`.
- `FileChannel.MapMode.READ_WRITE`: The resulting buffer is writable, and the changes will be written back to the file at some time. Note that other programs that have mapped the same file might not see those changes immediately. The exact behavior of simultaneous file mapping by multiple programs depends on the operating system.
- `FileChannel.MapMode.PRIVATE`: The resulting buffer is writable, but any changes are private to this buffer and not propagated to the file.

Once you have the buffer, you can read and write data using the methods of the `ByteBuffer` class and the `Buffer` superclass.

Buffers support both sequential and random data access. A buffer has a *position* that is advanced by `get` and `put` operations. For example, you can sequentially traverse all bytes in the buffer as

[Click here to view code image](#)

```
while (buffer.hasRemaining())
{
    byte b = buffer.get();
    . . .
}
```

Alternatively, you can use random access:

[Click here to view code image](#)

```
for (int i = 0; i < buffer.limit(); i++)
{
    byte b = buffer.get(i);
    . . .
}
```

You can also read and write arrays of bytes with the methods

[Click here to view code image](#)

```
get(byte[] bytes)
get(byte[], int offset, int length)
```

Finally, there are methods

[Click here to view code image](#)

```
getInt
getLong
```

```
getShort  
getChar  
getFloat  
getDouble
```

to read primitive type values that are stored as *binary* values in the file. As we already mentioned, Java uses big-endian ordering for binary data. However, if you need to process a file containing binary numbers in little-endian order, simply call

[Click here to view code image](#)

```
buffer.order(ByteOrder.LITTLE_ENDIAN);
```

To find out the current byte order of a buffer, call

```
ByteOrder b = buffer.order()
```



Caution

This pair of methods does not use the `set/get` naming convention.

To write numbers to a buffer, use one of the methods

[Click here to view code image](#)

```
putInt  
putLong  
putShort  
putChar  
putFloat  
putDouble
```

At some point, and certainly when the channel is closed, these changes are written back to the file.

[Listing 2.5](#) computes the 32-bit cyclic redundancy checksum (CRC32) of a file. That checksum is often used to determine whether a file has been corrupted. Corruption of a file makes it very likely that the checksum has changed. The `java.util.zip` package contains a class `CRC32` that computes the checksum of a sequence of bytes, using the following loop:

[Click here to view code image](#)

```
CRC32 crc = new CRC32();  
while (more bytes)  
    crc.update(next byte)  
long checksum = crc.getValue();
```



Note

For a nice explanation of the CRC algorithm, see www.relisoft.com/Science/CrcMath.html.

The details of the CRC computation are not important. We just use it as an example of a useful file operation. (In practice, you would read and update data in larger blocks, not a byte at a

time. Then the speed differences are not as dramatic.)

Run the program as

[Click here to view code image](#)

```
java memoryMap.MemoryMapTest filename
```

Listing 2.5 memoryMap/MemoryMapTest.java

[Click here to view code image](#)

```
1  package memoryMap;
2
3  import java.io.*;
4  import java.nio.*;
5  import java.nio.channels.*;
6  import java.nio.file.*;
7  import java.util.zip.*;
8
9  /**
10   * This program computes the CRC checksum of a file in four ways. <br>
11   * Usage: java memoryMap.MemoryMapTest filename
12   * @version 1.01 2012-05-30
13   * @author Cay Horstmann
14   */
15  public class MemoryMapTest
16  {
17      public static long checksumInputStream(Path filename) throws IOException
18      {
19          try (InputStream in = Files.newInputStream(filename))
20          {
21              CRC32 crc = new CRC32();
22
23              int c;
24              while ((c = in.read()) != -1)
25                  crc.update(c);
26              return crc.getValue();
27          }
28      }
29
30      public static long checksumBufferedInputStream(Path filename) throws
IOException
31      {
32          try (InputStream in = new
BufferedInputStream(Files.newInputStream(filename)))
33          {
34              CRC32 crc = new CRC32();
35
36              int c;
37              while ((c = in.read()) != -1)
38                  crc.update(c);
39              return crc.getValue();
40          }
41      }
42
43      public static long checksumRandomAccessFile(Path filename) throws
IOException
44      {
45          try (RandomAccessFile file = new RandomAccessFile(filename.toFile(),
```

```

"r"))
46     {
47         long length = file.length();
48         CRC32 crc = new CRC32();
49
50         for (long p = 0; p < length; p++)
51         {
52             file.seek(p);
53             int c = file.readByte();
54             crc.update(c);
55         }
56         return crc.getValue();
57     }
58 }
59
60 public static long checksumMappedFile(Path filename) throws IOException
61 {
62     try (FileChannel channel = FileChannel.open(filename))
63     {
64         CRC32 crc = new CRC32();
65         int length = (int) channel.size();
66         MappedByteBuffer buffer = channel.map(FileChannel.MapMode.READ_ONLY,
0, length);
67
68         for (int p = 0; p < length; p++)
69         {
70             int c = buffer.get(p);
71             crc.update(c);
72         }
73         return crc.getValue();
74     }
75 }
76
77 public static void main(String[] args) throws IOException
78 {
79     System.out.println("Input Stream:");
80     long start = System.currentTimeMillis();
81     Path filename = Paths.get(args[0]);
82     long crcValue = checksumInputStream(filename);
83     long end = System.currentTimeMillis();
84     System.out.println(Long.toHexString(crcValue));
85     System.out.println((end - start) + " milliseconds");
86
87     System.out.println("Buffered Input Stream:");
88     start = System.currentTimeMillis();
89     crcValue = checksumBufferedInputStream(filename);
90     end = System.currentTimeMillis();
91     System.out.println(Long.toHexString(crcValue));
92     System.out.println((end - start) + " milliseconds");
93
94     System.out.println("Random Access File:");
95     start = System.currentTimeMillis();
96     crcValue = checksumRandomAccessFile(filename);
97     end = System.currentTimeMillis();
98     System.out.println(Long.toHexString(crcValue));
99     System.out.println((end - start) + " milliseconds");
100
101     System.out.println("Mapped File:");
102     start = System.currentTimeMillis();
103     crcValue = checksumMappedFile(filename);

```

```

104         end = System.currentTimeMillis();
105         System.out.println(Long.toHexString(crcValue));
106         System.out.println((end - start) + " milliseconds");
107     }
108 }

```

java.io.FileInputStream 1.0

- `FileChannel getChannel()` **1.4**
returns a channel for accessing this input stream.
-
-

java.io.FileOutputStream 1.0

- `FileChannel getChannel()` **1.4**
returns a channel for accessing this output stream.
-
-

java.io.RandomAccessFile 1.0

- `FileChannel getChannel()` **1.4**
returns a channel for accessing this file.
-
-

java.nio.channels.FileChannel 1.4

- `static FileChannel open(Path path, OpenOption... options)` **7**
opens a file channel for the given path. By default, the channel is opened for reading.

<i>Parameters:</i>	<code>path</code>	The path to the file on which to open the channel
	<code>options</code>	Values <code>WRITE</code> , <code>APPEND</code> , <code>TRUNCATE_EXISTING</code> , <code>CREATE</code> in the <code>StandardOpenOption</code> enumeration

- `MappedByteBuffer map(FileChannel.MapMode mode, long position, long size)`
maps a region of the file to memory.

<i>Parameters:</i>	<code>mode</code>	One of the constants <code>READ_ONLY</code> , <code>READ_WRITE</code> , or <code>PRIVATE</code> in the <code>FileChannel.MapMode</code> class
	<code>position</code>	The start of the mapped region
	<code>size</code>	The size of the mapped region

java.nio.Buffer 1.4

- `boolean hasRemaining()`

returns `true` if the current buffer position has not yet reached the buffer's limit position.

- `int limit()`

returns the limit position of the buffer—that is, the first position at which no more values are available.

java.nio.ByteBuffer 1.4

- `byte get()`

gets a byte from the current position and advances the current position to the next byte.

- `byte get(int index)`

gets a byte from the specified index.

- `ByteBuffer put(byte b)`

puts a byte at the current position and advances the current position to the next byte. Returns a reference to this buffer.

- `ByteBuffer put(int index, byte b)`

puts a byte at the specified index. Returns a reference to this buffer.

- `ByteBuffer get(byte[] destination)`

- `ByteBuffer get(byte[] destination, int offset, int length)`

fills a byte array, or a region of a byte array, with bytes from the buffer, and advances the current position by the number of bytes read. If not enough bytes remain in the buffer, then no bytes are read, and a `BufferUnderflowException` is thrown. Returns a reference to this buffer.

<i>Parameters:</i>	<code>destination</code>	The byte array to be filled
	<code>offset</code>	The offset of the region to be filled
	<code>length</code>	The length of the region to be filled

- `ByteBuffer put(byte[] source)`

- `ByteBuffer put(byte[] source, int offset, int length)`

puts all bytes from a byte array, or the bytes from a region of a byte array, into the buffer, and advances the current position by the number of bytes read. If not enough bytes remain in the buffer, then no bytes are written, and a `BufferOverflowException` is thrown. Returns a reference to this buffer.

<i>Parameters:</i>	<code>source</code>	The byte array to be written
	<code>offset</code>	The offset of the region to be written
	<code>length</code>	The length of the region to be written

- `Xxx getXxx()`
- `Xxx getXxx(int index)`
- `ByteBuffer putXxx(Xxx value)`
- `ByteBuffer putXxx(int index, Xxx value)`

gets or puts a binary number. `Xxx` is one of `Int`, `Long`, `Short`, `Char`, `Float`, or `Double`.

- `ByteBuffer order(ByteOrder order)`
- `ByteOrder order()`

sets or gets the byte order. The value for `order` is one of the constants `BIG_ENDIAN` or `LITTLE_ENDIAN` of the `ByteOrder` class.

- `static ByteBuffer allocate(int capacity)`
constructs a buffer with the given capacity.
- `static ByteBuffer wrap(byte[] values)`
constructs a buffer that is backed by the given array.
- `CharBuffer asCharBuffer()`
constructs a character buffer that is backed by this buffer. Changes to the character buffer will show up in this buffer, but the character buffer has its own position, limit, and mark.

java.nio.CharBuffer 1.4

- `char get()`
- `CharBuffer get(char[] destination)`
- `CharBuffer get(char[] destination, int offset, int length)`

gets one `char` value, or a range of `char` values, starting at the buffer's position and moving the position past the characters that were read. The last two methods return this.

- `CharBuffer put(char c)`
- `CharBuffer put(char[] source)`
- `CharBuffer put(char[] source, int offset, int length)`
- `CharBuffer put(String source)`
- `CharBuffer put(CharBuffer source)`

puts one `char` value, or a range of `char` values, starting at the buffer's position and advancing the position past the characters that were written. When reading from a `CharBuffer`, all remaining characters are read. All methods return this.

2.6.2 The Buffer Data Structure

When you use memory mapping, you make a single buffer that spans the entire file or the area of the file that you're interested in. You can also use buffers to read and write more modest chunks of information.

In this section, we briefly describe the basic operations on `Buffer` objects. A buffer is an array of values of the same type. The `Buffer` class is an abstract class with concrete subclasses `ByteBuffer`, `CharBuffer`, `DoubleBuffer`, `FloatBuffer`, `IntBuffer`, `LongBuffer`, and `ShortBuffer`.



Note

The `StringBuffer` class is not related to these buffers.

In practice, you will most commonly use `ByteBuffer` and `CharBuffer`. As shown in [Figure 2.10](#), a buffer has

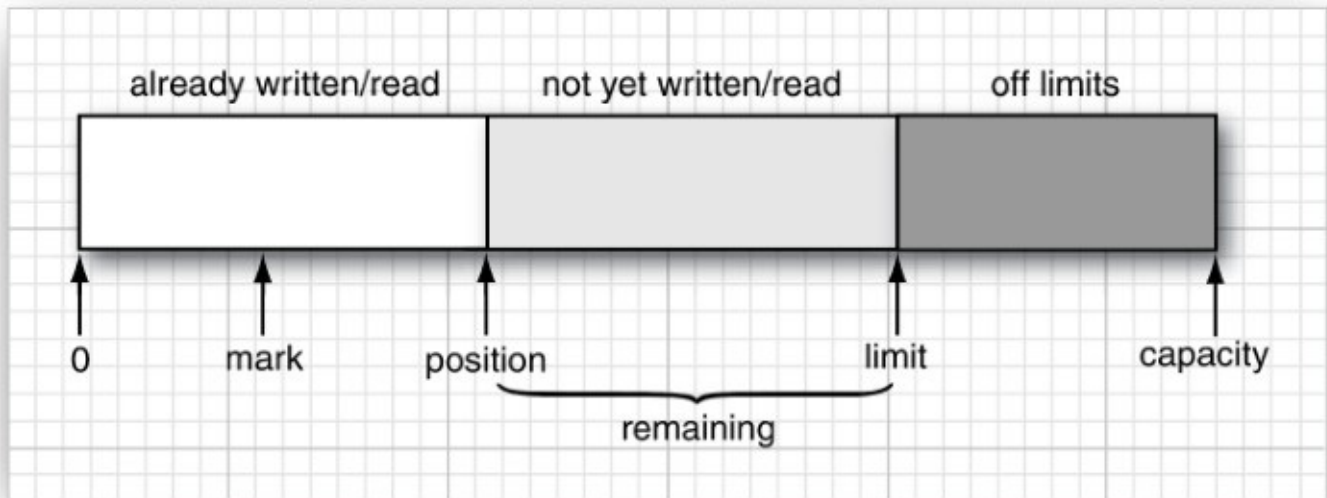


Figure 2.10 A buffer

- A *capacity* that never changes
- A *position* at which the next value is read or written
- A *limit* beyond which reading and writing is meaningless
- Optionally, a *mark* for repeating a read or write operation

These values fulfill the condition

$$0 = \text{mark} = \text{position} = \text{limit} = \text{capacity}$$

The principal purpose of a buffer is a “write, then read” cycle. At the outset, the buffer’s position is 0 and the limit is the capacity. Keep calling `put` to add values to the buffer. When you run out of data or reach the capacity, it is time to switch to reading.

Call `flip` to set the limit to the current position and the position to 0. Now keep calling `get`

while the `remaining` method (which returns *limit – position*) is positive. When you have read all values in the buffer, call `clear` to prepare the buffer for the next writing cycle. The `clear` method resets the position to 0 and the limit to the capacity.

If you want to reread the buffer, use `rewind` or `mark/reset` (see the API notes for details).

To get a buffer, call a static method such as `ByteBuffer.allocate` or `ByteBuffer.wrap`.

Then, you can fill a buffer from a channel, or write its contents to a channel. For example,

[Click here to view code image](#)

```
ByteBuffer buffer = ByteBuffer.allocate(RECORD_SIZE);
channel.read(buffer);
channel.position(newpos);
buffer.flip();
channel.write(buffer);
```

This can be a useful alternative to a random-access file.

java.nio.Buffer 1.4

- `Buffer clear()`
prepares this buffer for writing by setting the position to 0 and the limit to the capacity; returns `this`.
 - `Buffer flip()`
prepares this buffer for reading after writing, by setting the limit to the position and the position to 0; returns `this`.
 - `Buffer rewind()`
prepares this buffer for rereading the same values by setting the position to 0 and leaving the limit unchanged; returns `this`.
 - `Buffer mark()`
sets the mark of this buffer to the position; returns `this`.
 - `Buffer reset()`
sets the position of this buffer to the mark, thus allowing the marked portion to be read or written again; returns `this`.
 - `int remaining()`
returns the remaining number of readable or writable values—that is, the difference between the limit and position.
 - `int position()`
 - `void position(int newValue)`
gets and sets the position of this buffer.
 - `int capacity()`
returns the capacity of this buffer.
-

2.6.3 File Locking

When multiple simultaneously executing programs need to modify the same file, they need to communicate in some way, or the file can easily become damaged. File locks can solve this problem. A file lock controls access to a file or a range of bytes within a file.

Suppose your application saves a configuration file with user preferences. If a user invokes two instances of the application, it could happen that both of them want to write the configuration file at the same time. In that situation, the first instance should lock the file. When the second instance finds the file locked, it can decide to wait until the file is unlocked or simply skip the writing process.

To lock a file, call either the `lock` or `tryLock` methods of the `FileChannel` class.

[Click here to view code image](#)

```
FileChannel = FileChannel.open(path);  
FileLock lock = channel.lock();
```

or

[Click here to view code image](#)

```
FileLock lock = channel.tryLock();
```

The first call blocks until the lock becomes available. The second call returns immediately, either with the lock or with `null` if the lock is not available. The file remains locked until the channel is closed or the `release` method is invoked on the lock.

You can also lock a portion of the file with the call

[Click here to view code image](#)

```
FileLock lock(long start, long size, boolean shared)
```

or

[Click here to view code image](#)

```
FileLock tryLock(long start, long size, boolean shared)
```

The `shared` flag is `false` to lock the file for both reading and writing. It is `true` for a *shared* lock, which allows multiple processes to read from the file, while preventing any process from acquiring an exclusive lock. Not all operating systems support shared locks. You may get an exclusive lock even if you just asked for a shared one. Call the `isShared` method of the `FileLock` class to find out which kind you have.



Note

If you lock the tail portion of a file and the file subsequently grows beyond the locked portion, the additional area is not locked. To lock all bytes, use a size of `Long.MAX_VALUE`.

Be sure to unlock the lock when you are done. As always, this is best done with a `try-with-resources` statement:

[Click here to view code image](#)

```
try (FileLock lock = channel.lock())
{
    access the locked file or segment
}
```

Keep in mind that file locking is system-dependent. Here are some points to watch for:

- On some systems, file locking is merely *advisory*. If an application fails to get a lock, it may still write to a file that another application has currently locked.
- On some systems, you cannot simultaneously lock a file and map it into memory.
- File locks are held by the entire Java virtual machine. If two programs are launched by the same virtual machine (such as an applet or application launcher), they can't each acquire a lock on the same file. The `lock` and `tryLock` methods will throw an `OverlappingFileLockException` if the virtual machine already holds another overlapping lock on the same file.
- On some systems, closing a channel releases all locks on the underlying file held by the Java virtual machine. You should therefore avoid multiple channels on the same locked file.
- Locking files on a networked file system is highly system-dependent and should probably be avoided.

java.nio.channels.FileChannel 1.4

- `FileLock lock()`
acquires an exclusive lock on the entire file. This method blocks until the lock is acquired.
- `FileLock tryLock()`
acquires an exclusive lock on the entire file, or returns null if the lock cannot be acquired.
- `FileLock lock(long position, long size, boolean shared)`
- `FileLock tryLock(long position, long size, boolean shared)`
acquires a lock on a region of the file. The first method blocks until the lock is acquired, and the second method returns null if the lock cannot be acquired.

<i>Parameters:</i>	<code>position</code>	The start of the region to be locked
	<code>size</code>	The size of the region to be locked
	<code>shared</code>	true for a shared lock, false for an exclusive lock

java.nio.channels.FileLock 1.4

- `void close()` 1.7
releases this lock.
-

2.7 Regular Expressions

Regular expressions are used to specify string patterns. You can use regular expressions whenever you need to locate strings that match a particular pattern. For example, one of our sample programs locates all hyperlinks in an HTML file by looking for strings of the pattern `.`

Of course, when specifying a pattern, the `. . .` notation is not precise enough. You need to specify exactly what sequence of characters is a legal match, using a special syntax to describe a pattern.

Here is a simple example. The regular expression

```
[Jj]ava.+
```

matches any string of the following form:

- The first letter is a `J` or `j`.
- The next three letters are `ava`.
- The remainder of the string consists of one or more arbitrary characters.

For example, the string `"javanese"` matches this particular regular expression, but the string `"Core Java"` does not.

As you can see, you need to know a bit of syntax to understand the meaning of a regular expression. Fortunately, for most purposes, a few straightforward constructs are sufficient.

- A *character class* is a set of character alternatives, enclosed in brackets, such as `[Jj]`, `[0-9]`, `[A-Za-z]`, or `[^0-9]`. Here the `-` denotes a range (all characters whose Unicode values fall between the two bounds), and `^` denotes the complement (all characters except those specified).
- To include a `-` inside a character class, make it the first or last item. To include a `]`, make it the first item. To include a `^`, put it anywhere but the beginning. You only need to escape `[` and `\`.
- There are many predefined character classes such as `\d` (digits) or `\p{Sc}` (Unicode currency symbol). See [Tables 2.6](#) and [2.7](#).

Expression	Description	Example
Characters		
<i>c</i> , not one of . * + ? { () [\ ^ \$	The character <i>c</i>]
.	Any character except line terminators, or any character if the DOTALL flag is set	
\x{ <i>p</i> }	The Unicode code point with hex code <i>p</i>	\x{10546}
\uhhhh, \xhhh, \0o, \0oo, \0ooo	The UTF-16 code unit with the given hex or octal value	\uFEFF
\a, \e, \f, \n, \r, \t	Alert (\x{7}), escape (\x{1B}), form feed (\x{B}), newline (\x{A}), carriage return (\x{D}), tab (\x{9})	\n
\cc, where <i>c</i> is in [A-Z] or one of @ [\] ^ _ ?	The control character corresponding to the character <i>c</i>	\cH is a backspace (\x{8})
\c, where <i>c</i> is not in [A-Za-z0-9]	The character <i>c</i>	\\
\Q ... \E	Everything between the start and the end of the quotation	\Q(...) \E matches the string (...)
Character Classes		
[<i>C</i> ₁ <i>C</i> ₂ ...], where <i>C</i> _{<i>i</i>} are characters, ranges <i>c-d</i> , or character classes	Any of the characters represented by <i>C</i> ₁ , <i>C</i> ₂ , ...	[0-9+-]
[^ ...]	Complement of a character class	[^\d\s]
[... && ...]	Intersection of character classes	[\p{L}&&[^\A-Za-z]]
\p{ ... }, \P{ ... }	A predefined character class (see Table 2.7); its complement	\p{L} matches a Unicode letter, and so does \pL—you can omit braces around a single letter

<code>\d, \D</code>	Digits (<code>[0-9]</code> , or <code>\p{Digit}</code> when the <code>UNICODE_CHARACTER_CLASS</code> flag is set); the complement	<code>\d+</code> is a sequence of digits
<code>\w, \W</code>	Word characters (<code>[a-zA-Z0-9_]</code> , or Unicode word characters when the <code>UNICODE_CHARACTER_CLASS</code> flag is set); the complement	
<code>\s, \S</code>	Spaces (<code>[\n\r\t\f\x{B}]</code> , or <code>\p{IsWhite_Space}</code> when the <code>UNICODE_CHARACTER_CLASS</code> flag is set); the complement	<code>\s*, \S*</code> is a comma surrounded by optional white space
<code>\h, \v, \H, \V</code>	Horizontal whitespace, vertical whitespace, their complements	
Sequences and Alternatives		
<code>XY</code>	Any string from <code>X</code> , followed by any string from <code>Y</code>	<code>[1-9][0-9]*</code> is a positive number without leading zero
<code>X Y</code>	Any string from <code>X</code> or <code>Y</code>	<code>http ftp</code>
Grouping		
<code>(X)</code>	Captures the match of <code>X</code>	<code>'([\^']**)'</code> captures the quoted text
<code>\n</code>	The <i>n</i> th group	<code>(['"])\. *\1</code> matches <code>'Fred'</code> or <code>"Fred"</code> but not <code>"Fred'</code>
<code>(?<name>X)</code>	Captures the match of <code>X</code> with the given name	<code>'(?<id>[A-Za-z0-9]+)'</code> captures the match with name <code>id</code>
<code>\k<name></code>	The group with the given name	<code>\k<id></code> matches the group with name <code>id</code>
<code>(?:X)</code>	Use parentheses without capturing <code>X</code>	In <code>(?:http ftp)://(.*)</code> , the match after <code>://</code> is <code>\1</code>

$(?f_1f_2 \dots :X),$ $(?f_1 \dots -f_k \dots :X),$ with f_i in $[dimsu\backslash x]$	Matches, but does not capture, X with the given flags on or off (after $-$)	$(?i:jpe?g)$ is a case-insensitive match
Other $(? \dots)$	See the Pattern API documentation	
Quantifiers		
$X?$	Optional X	$\backslash+?$ is an optional $+$ sign
X^*, X^+	0 or more X , 1 or more X	$[1-9][0-9]^+$ is an integer \geq 10
$X\{n\}, X\{n, \}, X\{m, n\}$	n times X , at least n times X , between m and n times X	$[0-7]\{1,3\}$ are one to three octal digits
$Q?$, where Q is a quantified expression	Reluctant quantifier, attempting the shortest match before trying longer matches	$.*(<.+?>).*$ captures the shortest sequence enclosed in angle brackets
Q^+ , where Q is a quantified expression	Possessive quantifier, taking the longest match without backtracking	$'[\wedge]^+'$ matches strings enclosed in single quotes and fails quickly on strings without a closing quote
Boundary Matches		
$^, \$$	Beginning, end of input (or beginning, end of line in multiline mode)	$^Java\$$ matches the input or line <code>Java</code>
$\backslash A, \backslash Z, \backslash z$	Beginning of input, end of input, absolute end of input (unchanged in multiline mode)	
$\backslash b, \backslash B$	Word boundary, nonword boundary	$\backslash bJava\backslash b$ matches the word <code>Java</code>
$\backslash R$	A Unicode line break	
$\backslash G$	The end of the previous match	

Table 2.6 Regular Expression Syntax

Character Class Name	Explanation
<i>posixClass</i>	<i>posixClass</i> is one of Lower, Upper, Alpha, Digit, Alnum, Punct, Graph, Print, Cntrl, XDigit, Space, Blank, ASCII, interpreted as POSIX or Unicode class, depending on the UNICODE_CHARACTER_CLASS flag
<i>IsScript</i> , <i>sc=Script</i> , <i>script=Script</i>	A script accepted by <code>Character.UnicodeScript.forName</code>
<i>InBlock</i> , <i>blk=Block</i> , <i>block=Block</i>	A block accepted by <code>Character.UnicodeBlock.forName</code>
<i>Category</i> , <i>InCategory</i> , <i>gc=Category</i> , <i>general_category=Category</i>	A one- or two-letter name for a Unicode general category
<i>IsProperty</i>	<i>Property</i> is one of Alphabetic, Ideographic, Letter, Lowercase, Uppercase, Titlecase, Punctuation, Control, White_Space, Digit, Hex_Digit, Join_Control, Noncharacter_Code_Point, Assigned
<i>javaMethod</i>	Invokes the method <code>Character.isMethod</code> (must not be deprecated)

Table 2.7 Predefined Character Class Names Used with `\p`

- Most characters match themselves, such as the `ava` characters in the preceding example.
- The `.` symbol matches any character (except possibly line terminators, depending on flag settings).
- Use `\` as an escape character. For example, `\.` matches a period and `\\` matches a backslash.
- `^` and `$` match the beginning and end of a line, respectively.
- If *X* and *Y* are regular expressions, then *XY* means “any match for *X* followed by a match for *Y*”. *X | Y* means “any match for *X* or *Y*”.
- You can apply *quantifiers* *X*⁺ (1 or more), *X*^{*} (0 or more), and *X*[?] (0 or 1) to an expression *X*.
- By default, a quantifier matches the largest possible repetition that makes the overall match succeed. You can modify that behavior with suffixes `?` (reluctant, or stingy, match: match the smallest repetition count) and `+` (possessive, or greedy, match: match the largest count even if that makes the overall match fail).

For example, the string `cab` matches `[a-z]*ab` but not `[a-z]**+ab`. In the first case, the expression `[a-z]*` only matches the character `c`, so that the characters `ab` match the remainder of the pattern. But the greedy version `[a-z]**+` matches the characters `cab`, leaving the remainder of the pattern unmatched.

- You can use *groups* to define subexpressions. Enclose the groups in `()`, for example,

`([+-]?) ([0-9]+)`. You can then ask the pattern matcher to return the match of each group or to refer back to a group with `\n` where *n* is the group number, starting with `\1`.

For example, here is a somewhat complex but potentially useful regular expression that describes decimal or hexadecimal integers:

```
[+-]?[0-9]+|0[Xx][0-9A-Fa-f]+
```

Unfortunately, the regular expression syntax is not completely standardized between various programs and libraries; there is a consensus on the basic constructs but many maddening differences in the details. The Java regular expression classes use a syntax that is similar to, but not quite the same as, the one used in the Perl language. [Table 2.6](#) shows all constructs of the Java syntax. For more information on the regular expression syntax, consult the API documentation for the `Pattern` class or the book *Mastering Regular Expressions* by Jeffrey E. F. Friedl (O'Reilly and Associates, 2006).

The simplest use for a regular expression is to test whether a particular string matches it. Here is how you program that test in Java. First, construct a `Pattern` object from a string containing the regular expression. Then, get a `Matcher` object from the pattern and call its `matches` method:

[Click here to view code image](#)

```
Pattern pattern = Pattern.compile(patternString);
Matcher matcher = pattern.matcher(input);
if (matcher.matches()) . . .
```

The input of the matcher is an object of any class that implements the `CharSequence` interface, such as a `String`, `StringBuilder`, or `CharBuffer`.

When compiling the pattern, you can set one or more flags, for example:

[Click here to view code image](#)

```
Pattern pattern = Pattern.compile(expression,
    Pattern.CASE_INSENSITIVE + Pattern.UNICODE_CASE);
```

Or you can specify them inside the pattern:

[Click here to view code image](#)

```
String regex = "(?iU:expression)";
```

Here are the flags:

- `Pattern.CASE_INSENSITIVE` or `i`: Match characters independently of the letter case. By default, this flag takes only US ASCII characters into account.
- `Pattern.UNICODE_CASE` or `u`: When used in combination with `CASE_INSENSITIVE`, use Unicode letter case for matching.
- `Pattern.UNICODE_CHARACTER_CLASS` or `U`: Select Unicode character classes instead of POSIX. Implies `UNICODE_CASE`.
- `Pattern.MULTILINE` or `m`: Make `^` and `$` match the beginning and end of a line, not the entire input.
- `Pattern.UNIX_LINES` or `d`: Only `'\n'` is a line terminator when matching `^` and `$` in

multiline mode.

- `Pattern.DOTALL` or `s`: Make the `.` symbol match all characters, including line terminators.
- `Pattern.COMMENTS` or `x`: Whitespace and comments (from `#` to the end of a line) are ignored.
- `Pattern.LITERAL`: The pattern is taken literally and must be matched exactly, except possibly for letter case.
- `Pattern.CANON_EQ`: Take canonical equivalence of Unicode characters into account. For example, `u` followed by `¨` (diaeresis) matches `ü`.

The last two flags cannot be specified inside a regular expression.

If you want to match elements in a collection or stream, turn the pattern into a predicate:

[Click here to view code image](#)

```
Stream<String> strings = . . . ;
Stream<String> result = strings.filter(pattern.asPredicate());
```

The result contains all strings that match the regular expression.

If the regular expression contains groups, the `Matcher` object can reveal the group boundaries. The methods

[Click here to view code image](#)

```
int start(int groupIndex)
int end(int groupIndex)
```

yield the starting index and the past-the-end index of a particular group.

You can simply extract the matched string by calling

```
String group(int groupIndex)
```

Group 0 is the entire input; the group index for the first actual group is 1. Call the `groupCount` method to get the total group count. For named groups, use the methods

[Click here to view code image](#)

```
int start(String groupName)
int end(String groupName)
String group(String groupName)
```

Nested groups are ordered by the opening parentheses. For example, given the pattern

[Click here to view code image](#)

```
(( [1-9] | 1[0-2] ) : ([0-5] [0-9] )) [ap]m
```

and the input

```
11:59am
```

the matcher reports the following groups

Group Index	Start	End	String
0	0	7	11:59am
1	0	5	11:59
2	0	2	11
3	3	5	59

[Listing 2.6](#) prompts for a pattern, then for strings to match. It prints out whether or not the input matches the pattern. If the input matches and the pattern contains groups, the program prints the group boundaries as parentheses, for example:

```
((11):(59))am
```

Listing 2.6 `regex/RegexTest.java`

[Click here to view code image](#)

```

1  package regex;
2
3  import java.util.*;
4  import java.util.regex.*;
5
6  /**
7   * This program tests regular expression matching. Enter a pattern and
strings to match,
8   * or hit Cancel to exit. If the pattern contains groups, the group
boundaries are displayed
9   * in the match.
10  * @version 1.02 2012-06-02
11  * @author Cay Horstmann
12  */
13  public class RegexTest
14  {
15      public static void main(String[] args) throws PatternSyntaxException
16      {
17          Scanner in = new Scanner(System.in);
18          System.out.println("Enter pattern: ");
19          String patternString = in.nextLine();
20
21          Pattern pattern = Pattern.compile(patternString);
22
23          while (true)
24          {
25              System.out.println("Enter string to match: ");
26              String input = in.nextLine();
27              if (input == null || input.equals("")) return;
28              Matcher matcher = pattern.matcher(input);
29              if (matcher.matches())
30              {
31                  System.out.println("Match ");
32                  int g = matcher.groupCount();
33                  if (g > 0)
34                  {
35                      for (int i = 0; i < input.length(); i++)
36                      {
37                          // Print any empty groups

```

```

38         for (int j = 1; j <= g; j++)
39             if (i == matcher.start(j) && i == matcher.end(j))
40                 System.out.print("(");
41         // Print ( for non-empty groups starting here
42         for (int j = 1; j <= g; j++)
43             if (i == matcher.start(j) && i != matcher.end(j))
44                 System.out.print('(');
45         System.out.print(input.charAt(i));
46         // Print ) for non-empty groups ending here
47         for (int j = 1; j <= g; j++)
48             if (i + 1 != matcher.start(j) && i + 1 ==
matcher.end(j))
49                 System.out.print(')');
50         }
51         System.out.println();
52     }
53 }
54 else
55     System.out.println("No match");
56 }
57 }
58 }

```

Usually, you don't want to match the entire input against a regular expression, but to find one or more matching substrings in the input. Use the `find` method of the `Matcher` class to find the next match. If it returns `true`, use the `start` and `end` methods to find the extent of the match or the `group` method without an argument to get the matched string.

[Click here to view code image](#)

```

while (matcher.find())
{
    int start = matcher.start();
    int end = matcher.end();
    String match = input.group();
    . . .
}

```

[Listing 2.7](#) puts this mechanism to work. It locates all hypertext references in a web page and prints them. To run the program, supply a URL on the command line, such as

[Click here to view code image](#)

```
java match.HrefMatch http://horstmann.com
```

Listing 2.7 `match/HrefMatch.java`

[Click here to view code image](#)

```

1  package match;
2
3  import java.io.*;
4  import java.net.*;
5  import java.nio.charset.*;
6  import java.util.regex.*;
7
8  /**
9   * This program displays all URLs in a web page by matching a regular
expression that describes

```

```

10  * the <a href=...> HTML tag. Start the program as <br>
11  * java match.HrefMatch URL
12  * @version 1.02 2016-07-14
13  * @author Cay Horstmann
14  */
15  public class HrefMatch
16  {
17      public static void main(String[] args)
18      {
19          try
20          {
21              // get URL string from command line or use default
22              String urlString;
23              if (args.length > 0) urlString = args[0];
24              else urlString = "http://java.sun.com";
25
26              // open reader for URL
27              InputStreamReader in = new InputStreamReader(new
URL(urlString).openStream(),
28                  StandardCharsets.UTF_8);
29
30              // read contents into string builder
31              StringBuilder input = new StringBuilder();
32              int ch;
33              while ((ch = in.read()) != -1)
34                  input.append((char) ch);
35
36              // search for all occurrences of pattern
37              String patternString = "<a\\s+href\\s*=\\s*(\\\"[^\\\"]*\\\"|
[^\s>]*)\\s*>";
38              Pattern pattern = Pattern.compile(patternString,
Pattern.CASE_INSENSITIVE);
39              Matcher matcher = pattern.matcher(input);
40
41              while (matcher.find())
42              {
43                  String match = matcher.group();
44                  System.out.println(match);
45              }
46          }
47          catch (IOException | PatternSyntaxException e)
48          {
49              e.printStackTrace();
50          }
51      }
52  }

```

The `replaceAll` method of the `Matcher` class replaces all occurrences of a regular expression with a replacement string. For example, the following instructions replace all sequences of digits with a `#` character:

[Click here to view code image](#)

```

Pattern pattern = Pattern.compile("[0-9]+");
Matcher matcher = pattern.matcher(input);
String output = matcher.replaceAll("#");

```

The replacement string can contain references to the groups in the pattern: `$n` is replaced with the *n*th group, and `${name}` is replaced with the group that has the given name. Use `\$` to

include a `$` character in the replacement text.

If you have a string that may contain `$` and `\`, and you don't want them to be interpreted as group replacements, call `matcher.replaceAll(Matcher.quoteReplacement(str))`.

The `replaceFirst` method replaces only the first occurrence of the pattern.

Finally, the `Pattern` class has a `split` method that splits an input into an array of strings, using the regular expression matches as boundaries. For example, the following instructions split the input into tokens, where the delimiters are punctuation marks surrounded by optional whitespace.

[Click here to view code image](#)

```
Pattern pattern = Pattern.compile("\\s*\\p{Punct}\\s*");
String[] tokens = pattern.split(input);
```

If there are many tokens, you can fetch them lazily:

[Click here to view code image](#)

```
Stream<String> tokens = commas.splitAsStream(input);
```

If you don't care about precompiling the pattern or lazy fetching, you can just use the `String.split` method:

[Click here to view code image](#)

```
String[] tokens = input.split("\\s*,\\s*");
```

java.util.regex.Pattern 1.4

- static Pattern compile(String expression)
- static Pattern compile(String expression, int flags)

compiles the regular expression string into a pattern object for fast processing of matches.

Parameters:

expression	The regular expression
flags	One or more of the flags CASE_INSENSITIVE, UNICODE_CASE, MULTILINE, UNIX_LINES, DOTALL, and CANON_EQ

- Matcher matcher(CharSequence input)

returns a matcher object that you can use to locate the matches of the pattern in the input.

- String[] split(CharSequence input)
- String[] split(CharSequence input, int limit)
- Stream<String> splitAsStream(CharSequence input) 8

splits the input string into tokens, where the pattern specifies the form of the delimiters. Returns an array or stream of tokens. The delimiters are not part of the tokens.

Parameters:

input	The string to be split into tokens
limit	The maximum number of strings to produce. If limit - 1 matching delimiters have been found, then the last entry of the returned array contains the remaining unsplit input. If limit is ≤ 0 , then the entire input is split. If limit is 0, then trailing empty strings are not placed in the returned array.

java.util.regex.Matcher 1.4

- boolean matches()
returns true if the input matches the pattern.
- boolean lookingAt()
returns true if the beginning of the input matches the pattern.
- boolean find()
- boolean find(int start)
attempts to find the next match and returns true if another match is found.

Parameters:

start	The index at which to start searching
-------	---------------------------------------

- int start()
- int end()
returns the start or past-the-end position of the current match.

- `String group()`
returns the current match.
 - `int groupCount()`
returns the number of groups in the input pattern.
 - `int start(int groupIndex)`
 - `int end(int groupIndex)`
returns the start or past-the-end position of a given group in the current match.
Parameters: `groupIndex` The group index (starting with 1), or 0 to indicate the entire match
 - `String group(int groupIndex)`
returns the string matching a given group.
Parameters: `groupIndex` The group index (starting with 1), or 0 to indicate the entire match
 - `String replaceAll(String replacement)`
 - `String replaceFirst(String replacement)`
returns a string obtained from the matcher input by replacing all matches, or the first match, with the replacement string.
Parameters: `replacement` The replacement string. It can contain references to pattern groups as `$n`. Use `\$` to include a `$` symbol.
 - `static String quoteReplacement(String str) 5.0`
quotes all `\` and `$` in `str`.
 - `Matcher reset()`
 - `Matcher reset(CharSequence input)`
resets the matcher state. The second method makes the matcher work on a different input. Both methods return this.
-

You have now seen how to carry out input and output operations in Java, and had an overview of the regular expression package that was a part of the “new I/O” specification. In the next chapter, we turn to the processing of XML data.